

Organiza



Cloud native monitoring with Prometheus

Beatriz Martínez | Cloud Engineer en IBM

 @beatrizmrg

IBMer, insatiable learner, passionate about technology and innovation #AI #cloud #devops



DEVOPS
SPAIN
II EDICIÓN

Patrorna



Colabora



DEVOPS SPAIN

II EDICIÓN



Talk overview

Cloud native monitoring
with Prometheus

Índice

1. Prometheus

Monitoring from the way we used to, to cloud-native world

2. Monitoring systems and applications

Metrics scraping, querying, graphing, instrumenting prometheus-native apps.

3. Alerting and Integrations

Decoupling, grouping, silencing, alerting.

4. Scaling and ecosystem

Federation, Thanos, k8s operator, Signifai



1. Prometheus

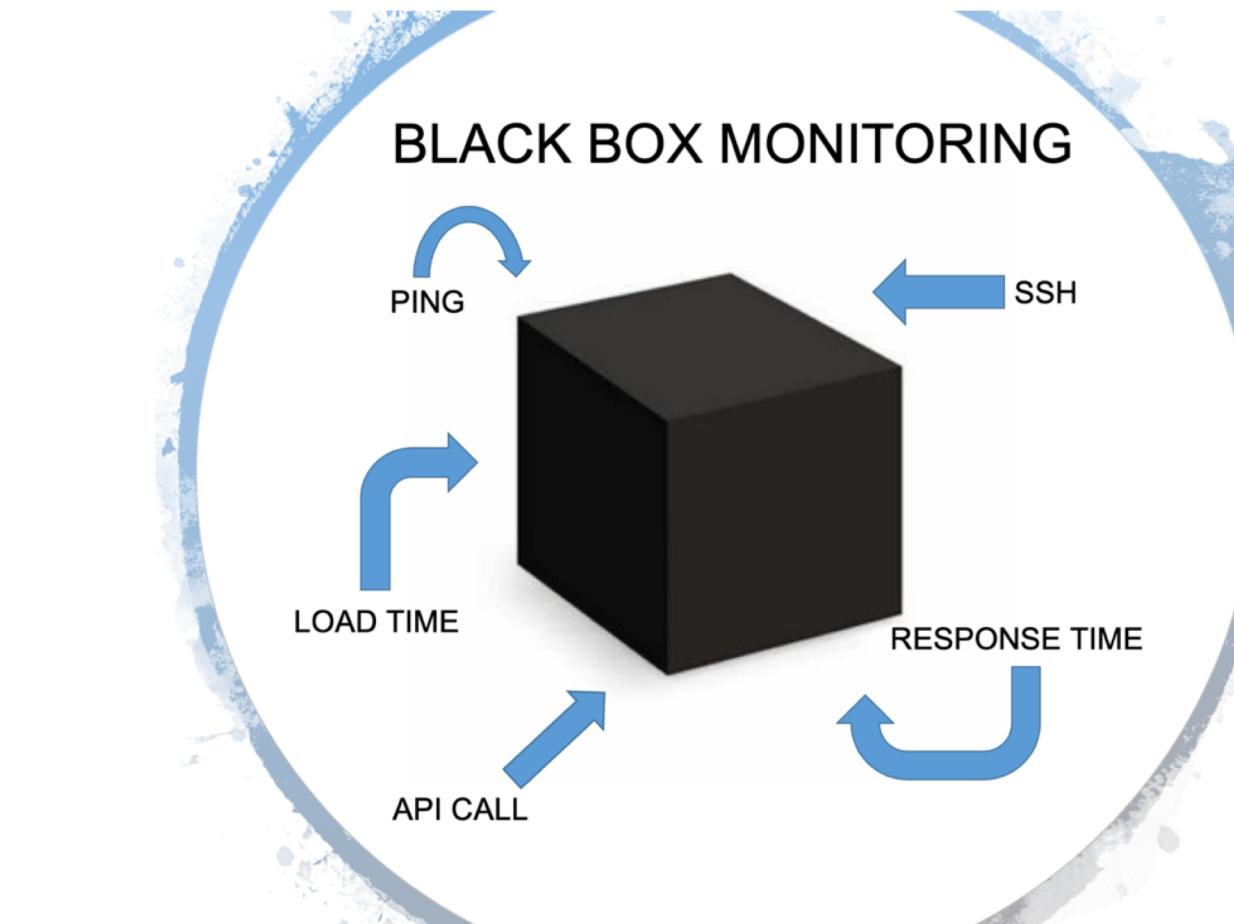
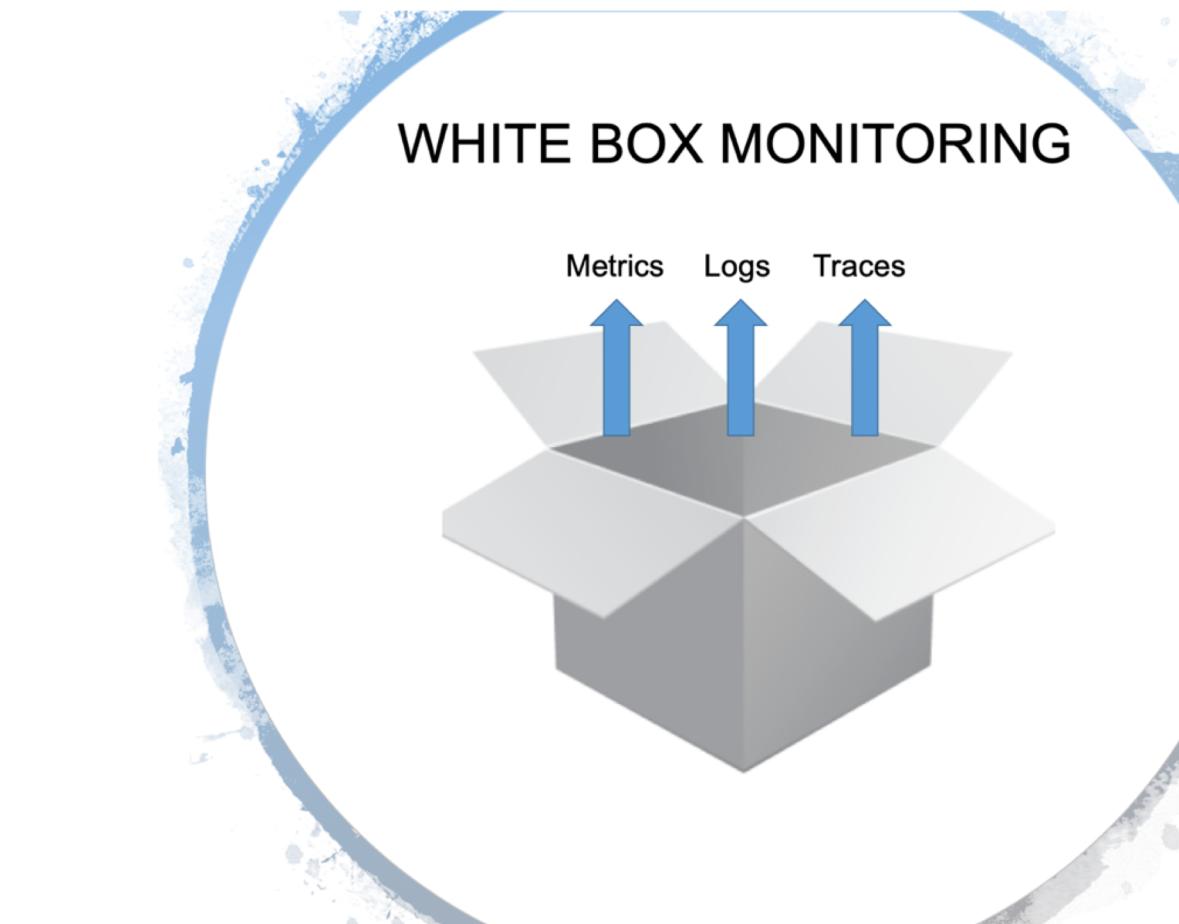
Monitoring from the way we used to, to cloud-native world

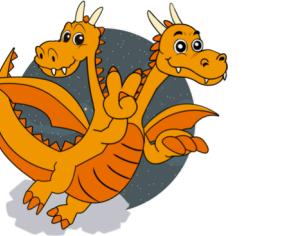
1.1 Monitoring

What do we understand for monitoring?

Ensure both systems and applications works as they should.

- Automatic or manual monitoring systems.
- Blackbox or Whitebox systems.



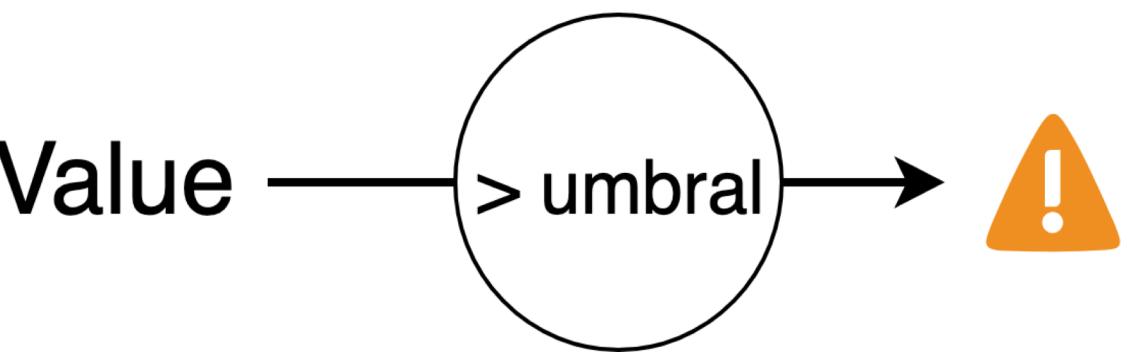


1.2 Monitoring systems space

Well known monitoring systems

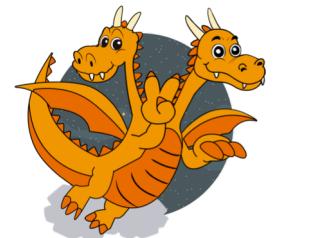
We used to ensure the system health based on some basic checks available such as:

- CPU load
- Memory use
- Disk use
- Network use



Monitoring systems with this behaviour:

- Nagios (1999)
- Sensu



1.2 Monitoring systems space

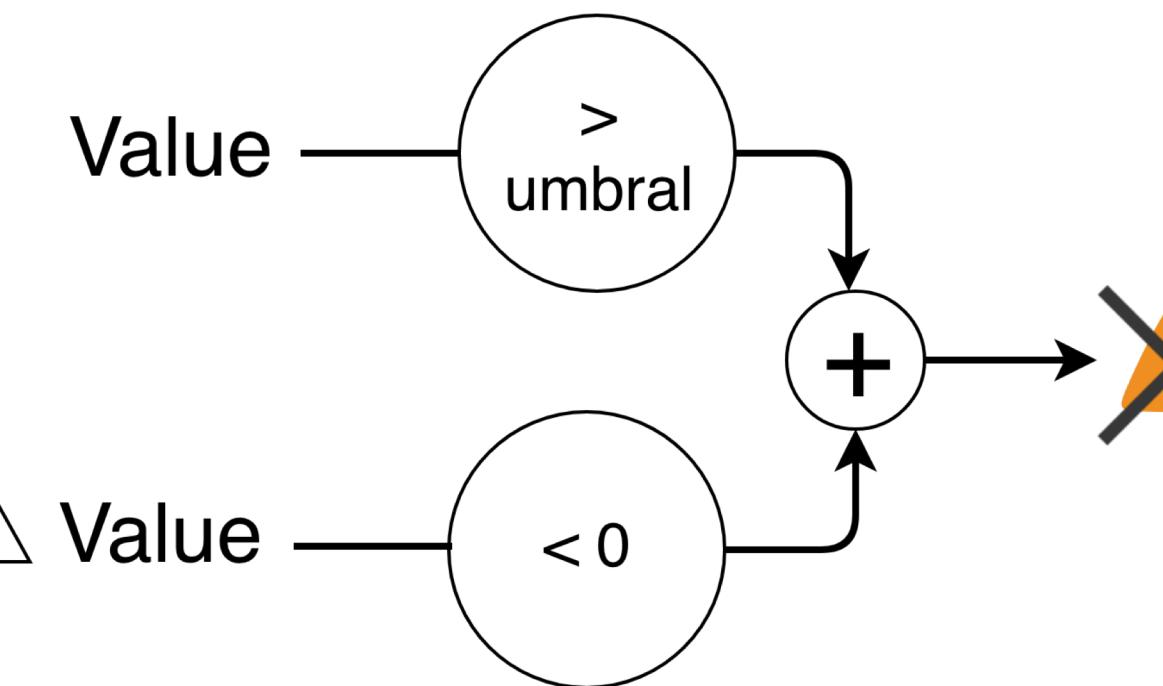
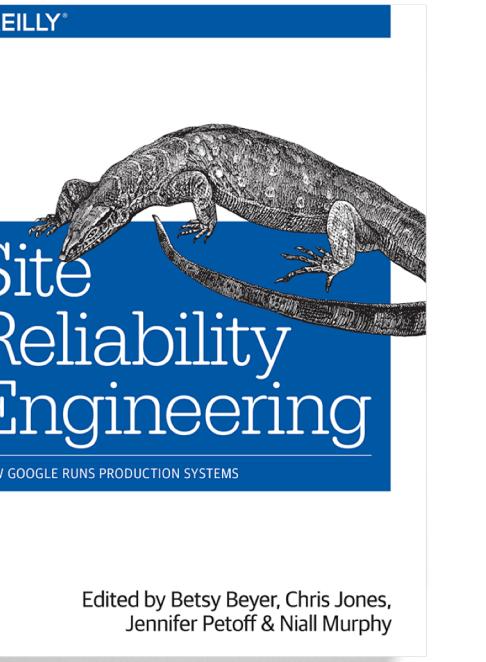
Today necessities

Today must-have metrics we should be obtaining out of any system (RED method):

- Request rate
- Error rate
- Duration, response time (Latency)

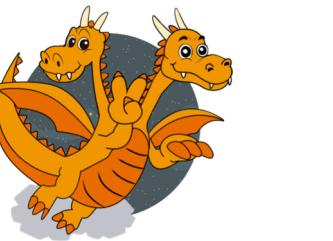
Monitoring systems metrics based:

- Graphite
- Prometheus



Allow alerts based on more complex queries and correlate data:

High CPU System A	High Error Rate System B
High Duration	High Request Rate



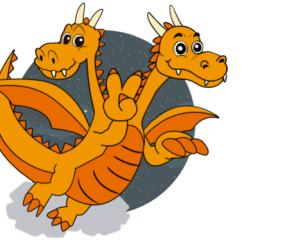
1.2 Monitoring systems space

Observability

- know when **things go wrong**
- Be able to **debug** and gain **insight**

- Trending to see changes over time, and **drive technical/business decisions**
- **Feed into other systems** (QA, security, automation, ...)





1.2 Monitoring systems space

Working with metrics

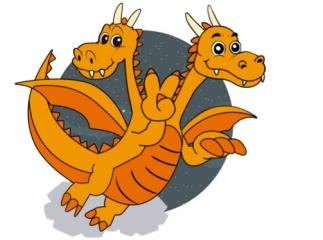
In order to get metrics of the system or app insight:

- **Logs** based metrics
- Specific **metrics** systems

Even being able to graph same events with both, they **complement** and are used for specific goals.

Logs always have higher **computational cost**, so correlations and complex queries made with metrics would be preferable.

```
x.x.x.90 - - [13/Sep/2006:07:01:53 -0700] "PROPFIND /svn/[xxxx]/Extranet/branches/SOW-101 HTTP/1.1" 401 587
x.x.x.90 - - [13/Sep/2006:07:01:51 -0700] "PROPFIND /svn/[xxxx]/[xxxx]/trunk HTTP/1.1" 401 587
x.x.x.90 - - [13/Sep/2006:07:00:53 -0700] "PROPFIND /svn/[xxxx]/[xxxx]/2.5 HTTP/1.1" 401 587
x.x.x.90 - - [13/Sep/2006:07:00:53 -0700] "PROPFIND /svn/[xxxx]/Extranet/branches/SOW-101 HTTP/1.1" 401 587
```



1.3 Prometheus

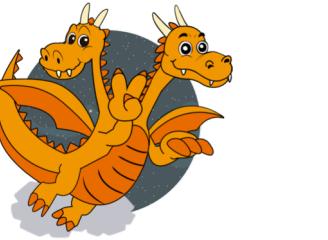
The Development History of this instrumentation and monitoring stack.



Borg —————→ Kubernetes
Borgmon —————→ Prometheus



- Inspired by Google's Borgmon monitoring system.
- **Started in 2012** by ex-Googlers working in Soundcloud as an **open source** project
- Mainly written in **Go**.
- Publicly **launched** in early **2015**
- Today continues to be independent of any company.
- K8s and Prometheus, along with envoy are the only three **graduated CNCF projects**.

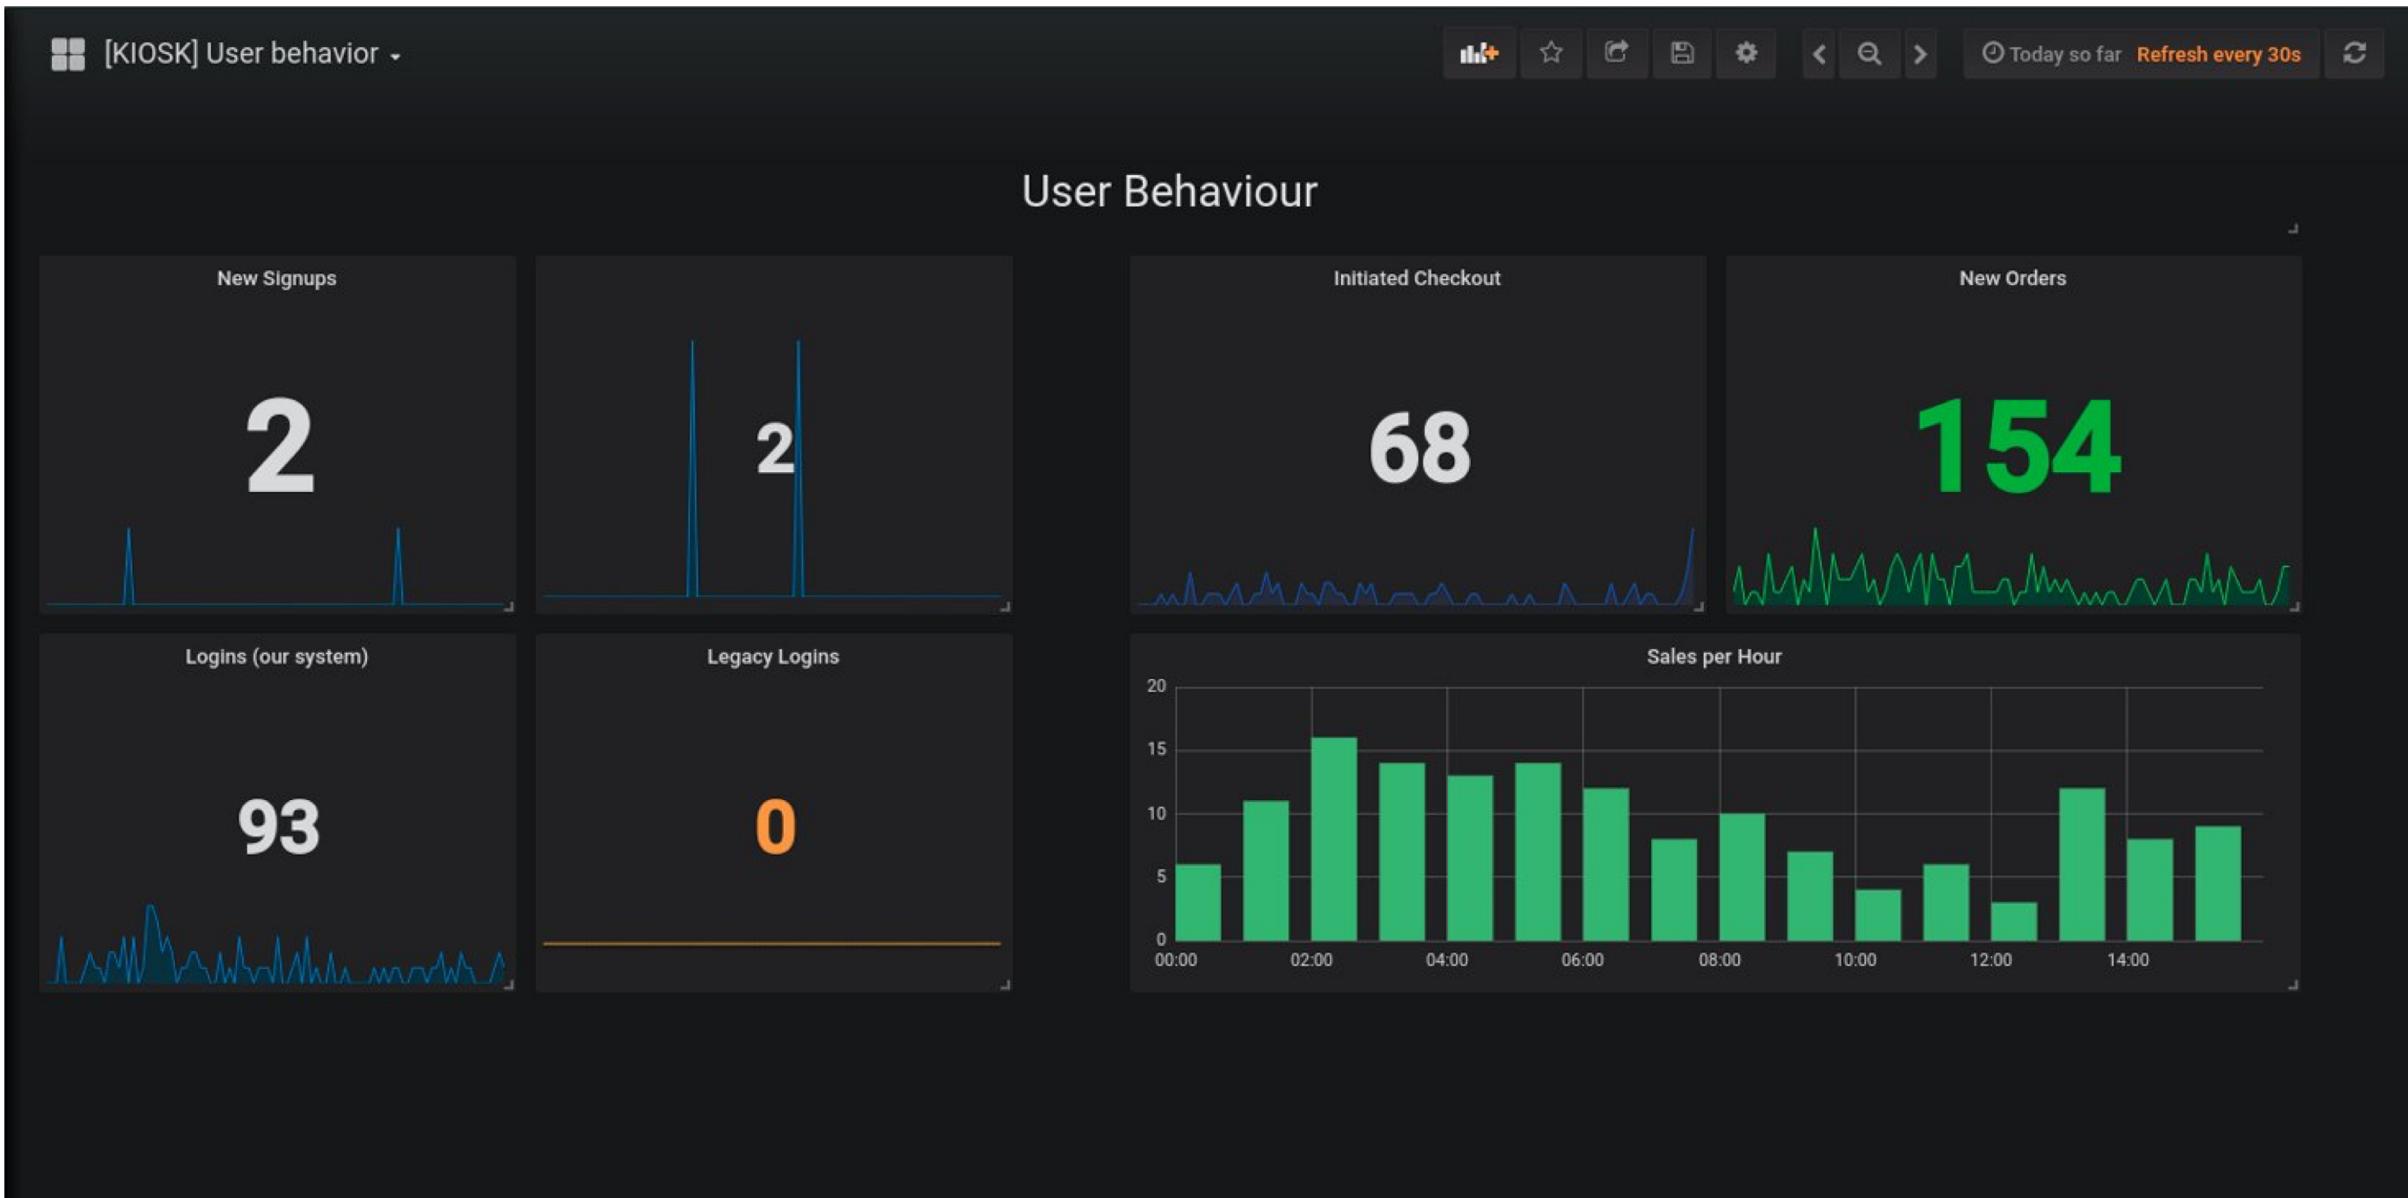


1.3 Prometheus

Why Prometheus?

It has knowledge about what the world should look like (which endpoints should exist, what time series means trouble, ..) and actively tries to find faults.

- Whitebox: Instrumenting applications and systems.
- Metrics based: **gaining insights** both systems and **business**.
- Manageable and Reliable
- Advanced service discovery
- Efficient: single server can handle millions of metrics, hundreds of thousands of datapoints per second
- Scalable
- Easy to integrate with
-



3 Prometheus

coming the standard

Prometheus metrics:

- Databases
 - Hardware related
 - Messaging systems
 - Storage
 - HTTP
 - APIs
 - Logging
 -

Global exporters:



Native:



1.3 Prometheus

Becoming the standard

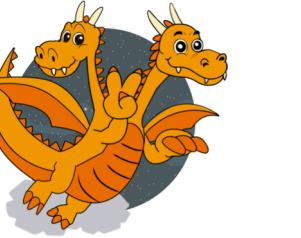
Open Metrics Project

- Standard based on prometheus exposition format.
- Exploit metrics beyond Prometheus.
- Accepted by **Cloud Native Computing Foundation**. Sandbox (early-stage project) Mars 2018.
- Aims to bring useful metrics to cloud-native deployments
- At its core, is an effort to develop a neutral metrics exposition format.
- Define an **efficient way to transport the metrics over the wire**, and a flexible way to attach information to them: label sets.



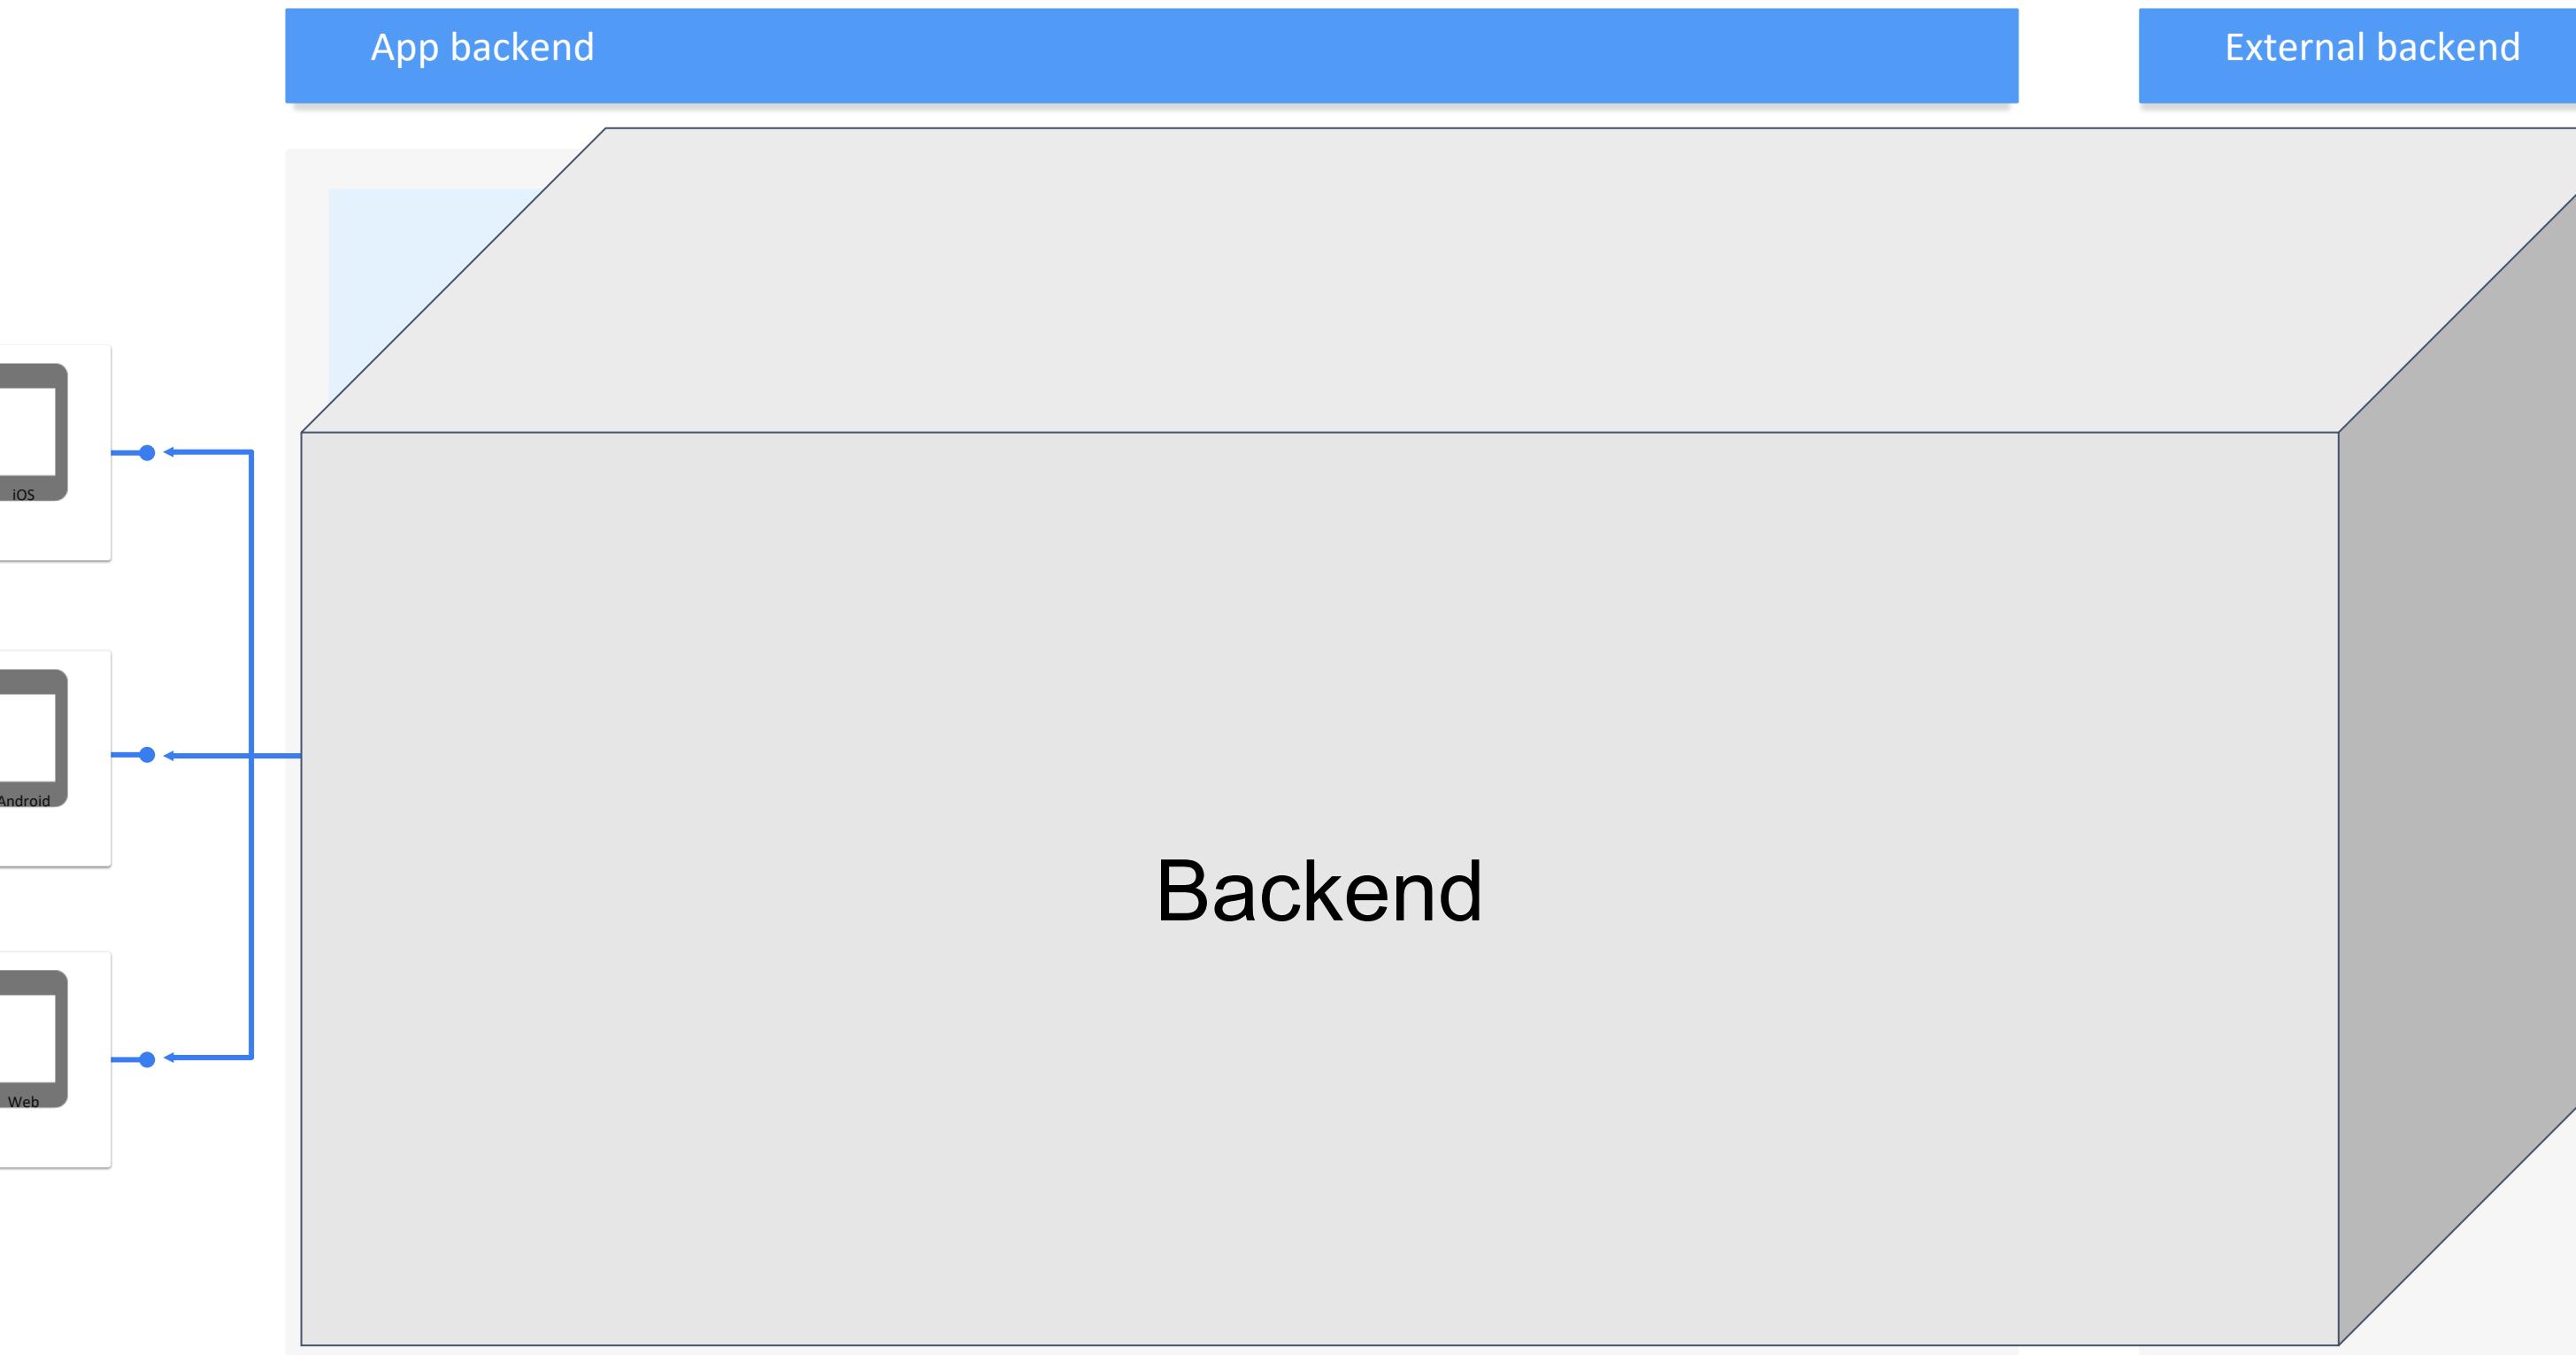
OPEN METRICS

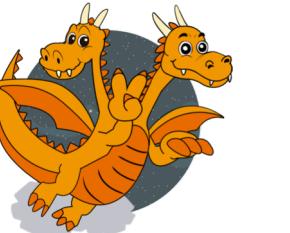
An effort to create an open standard for transmitting metrics at scale, with support for both **text representation** and **Protocol buffers**.



1.4 Cloud-native monitoring

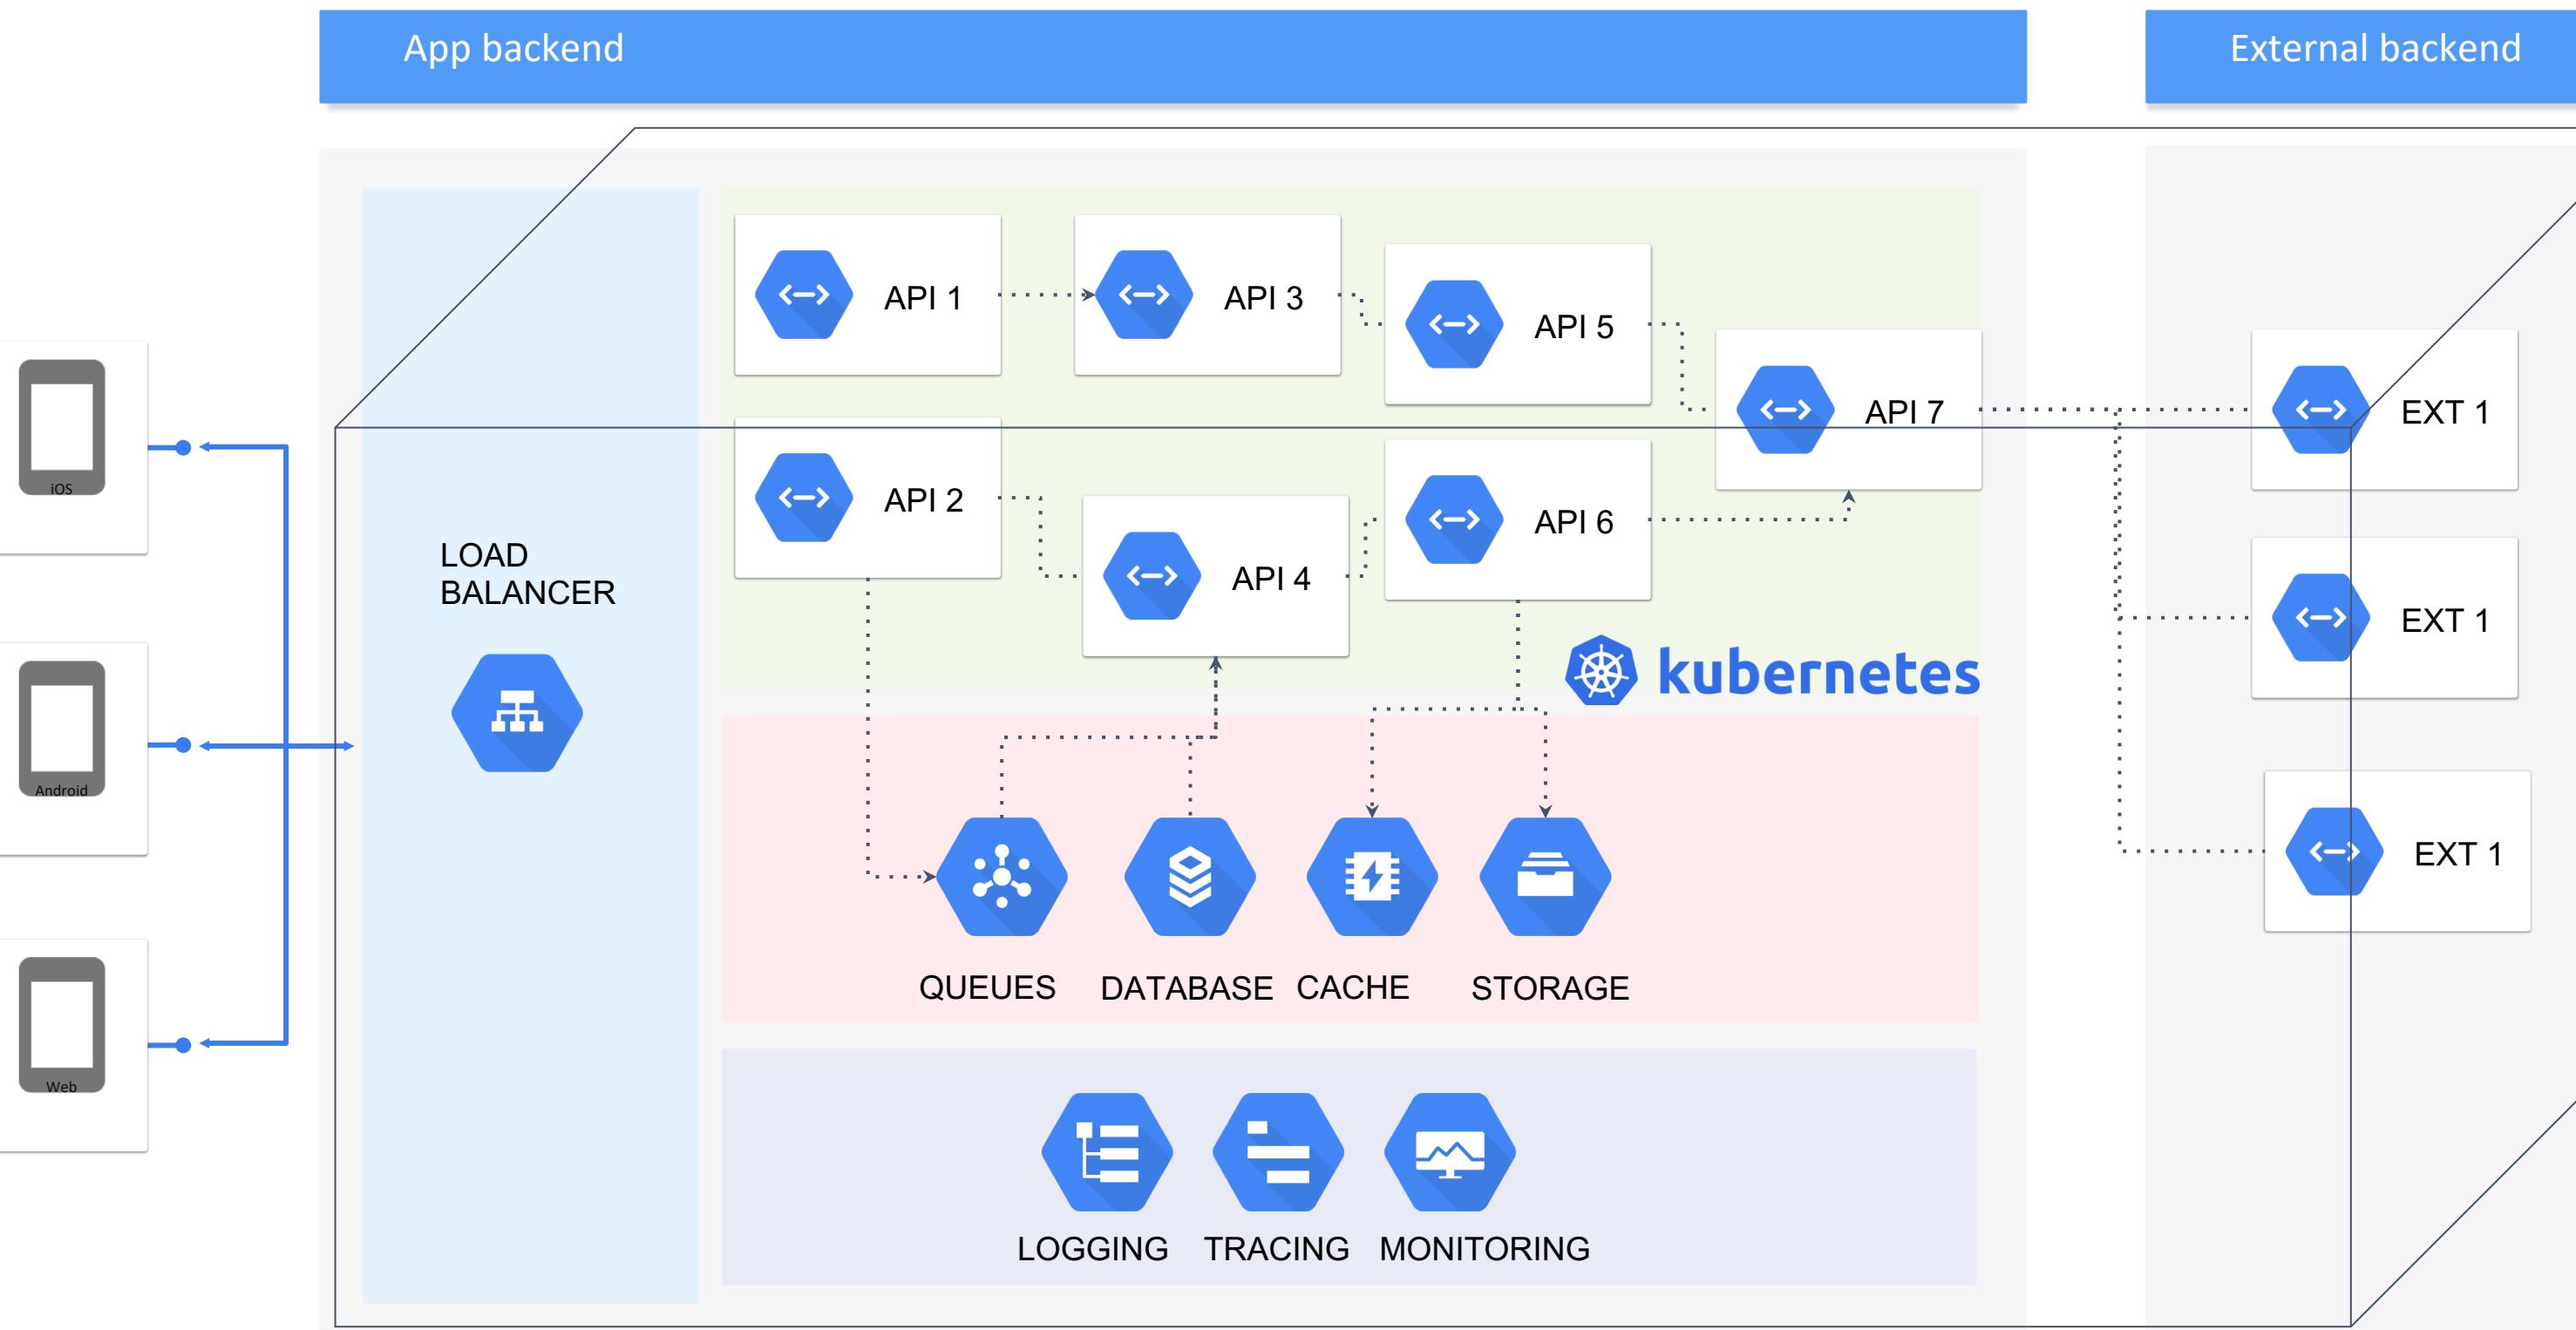
Microservices observability

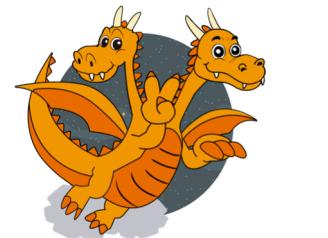




1.4 Cloud-native monitoring

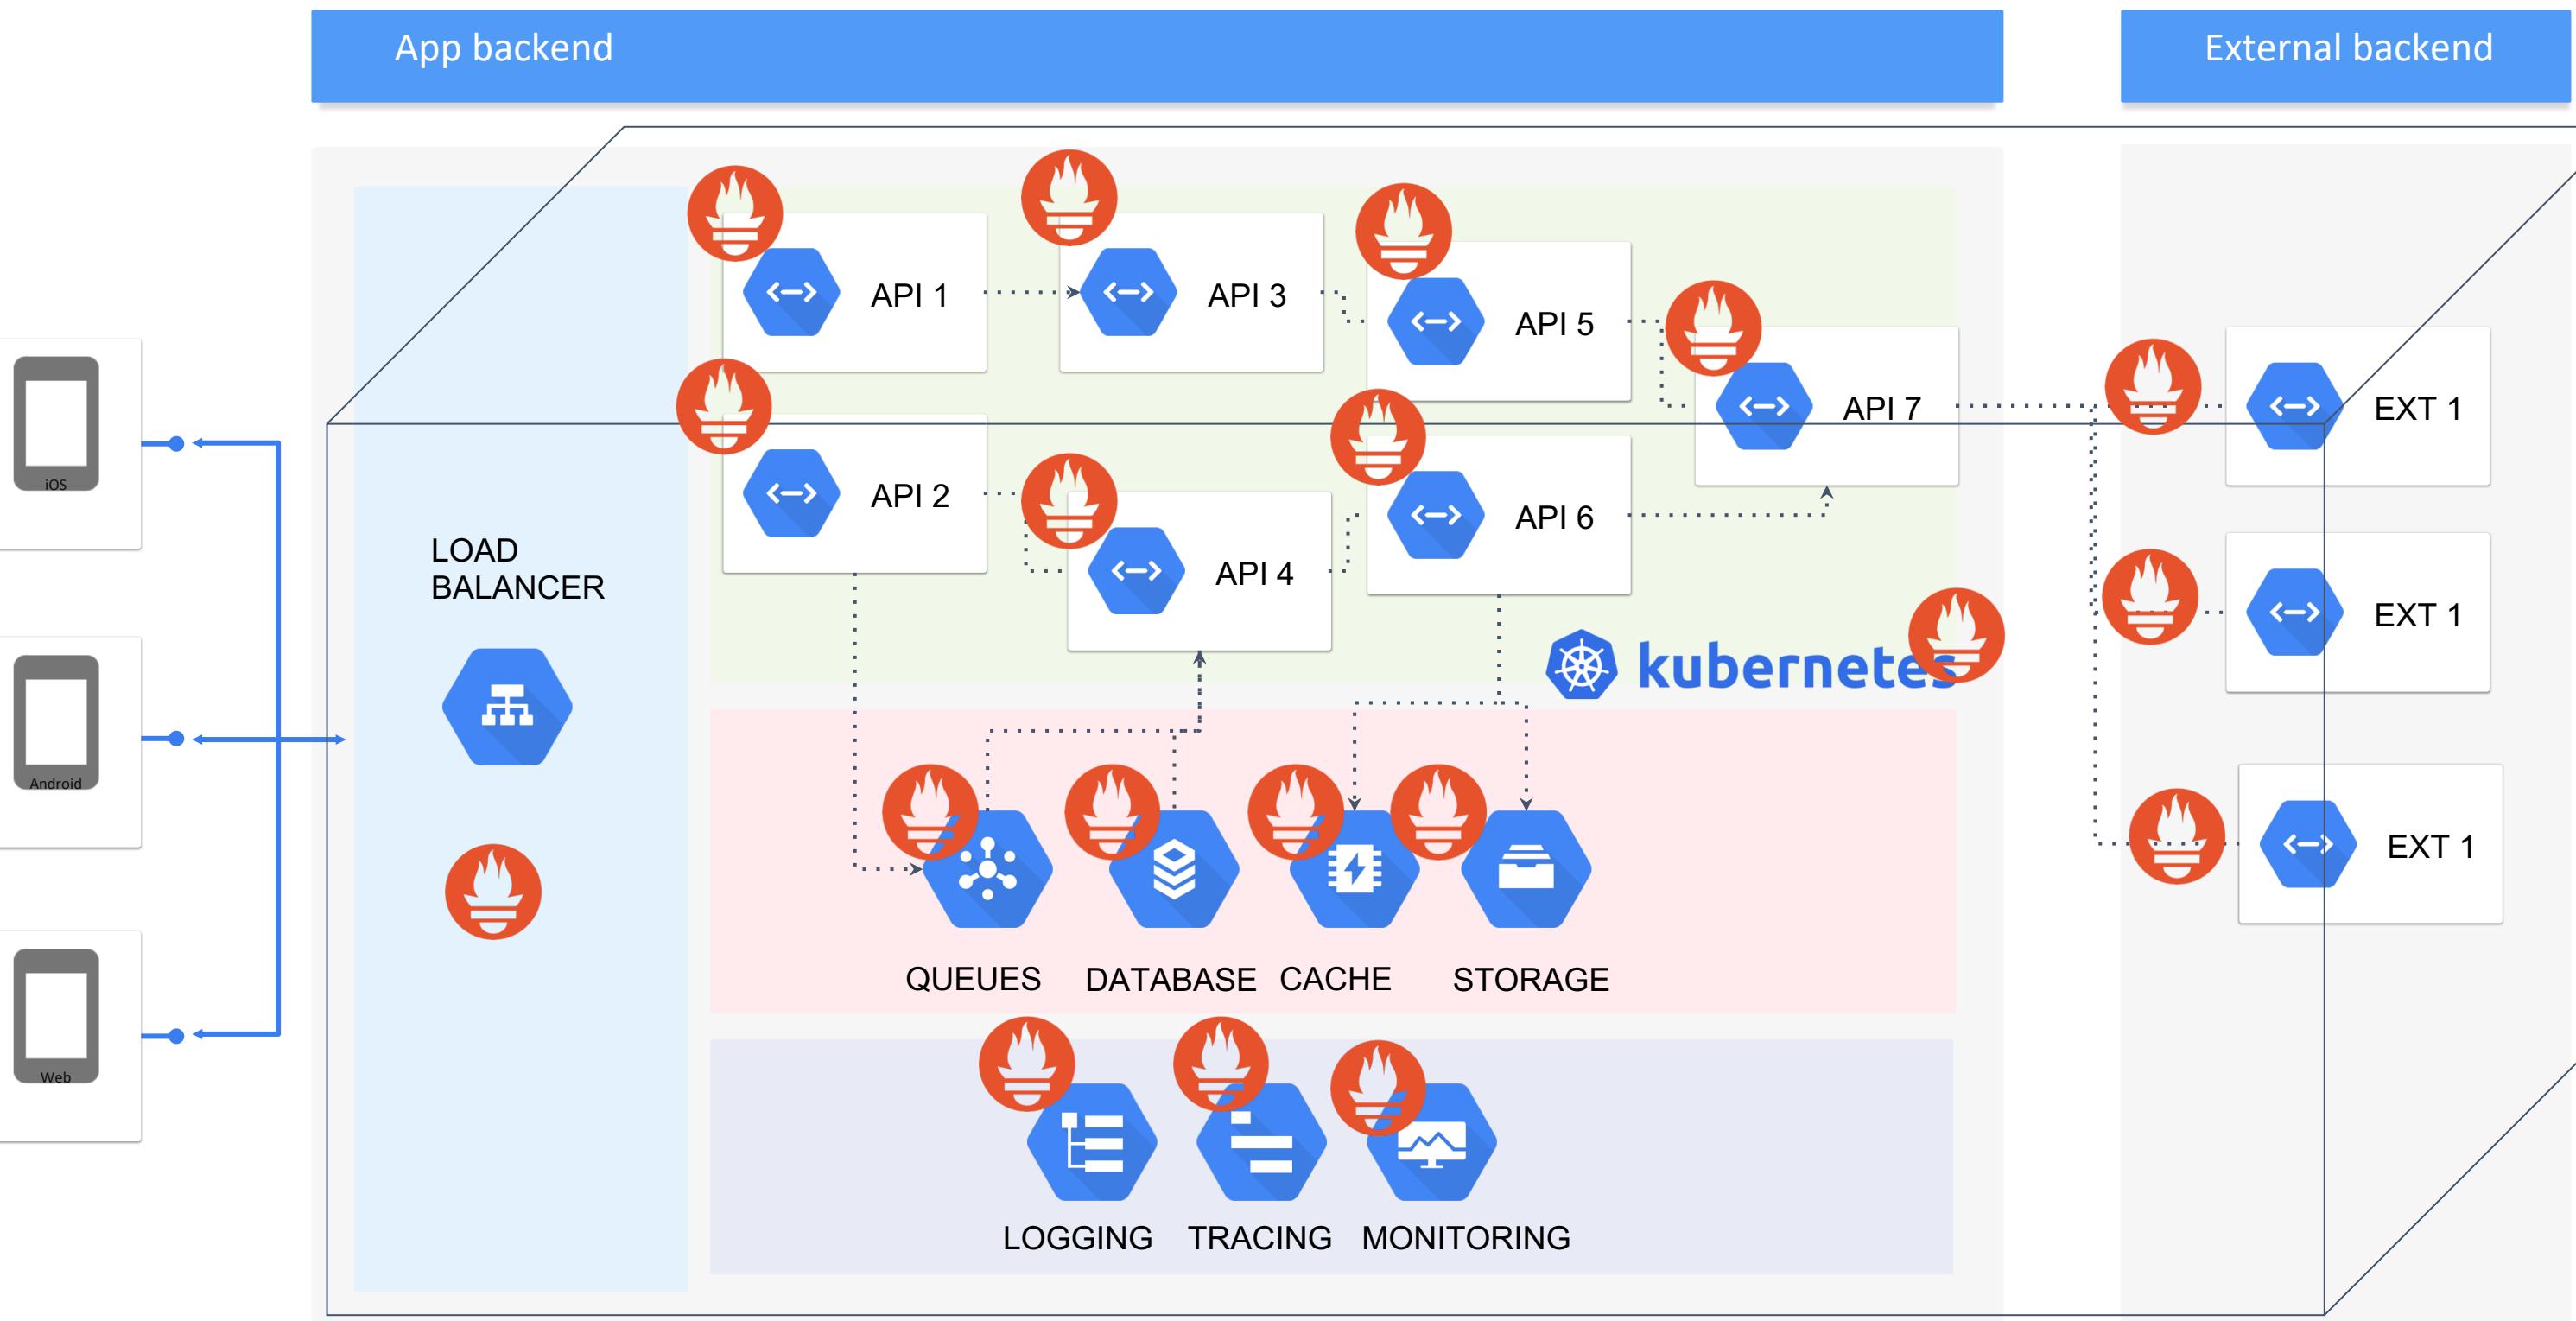
Microservices observability





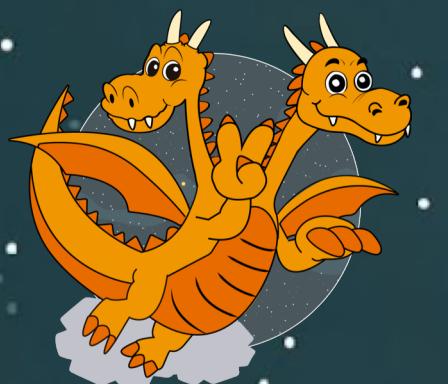
1.4 Cloud-native monitoring

Microservices observability



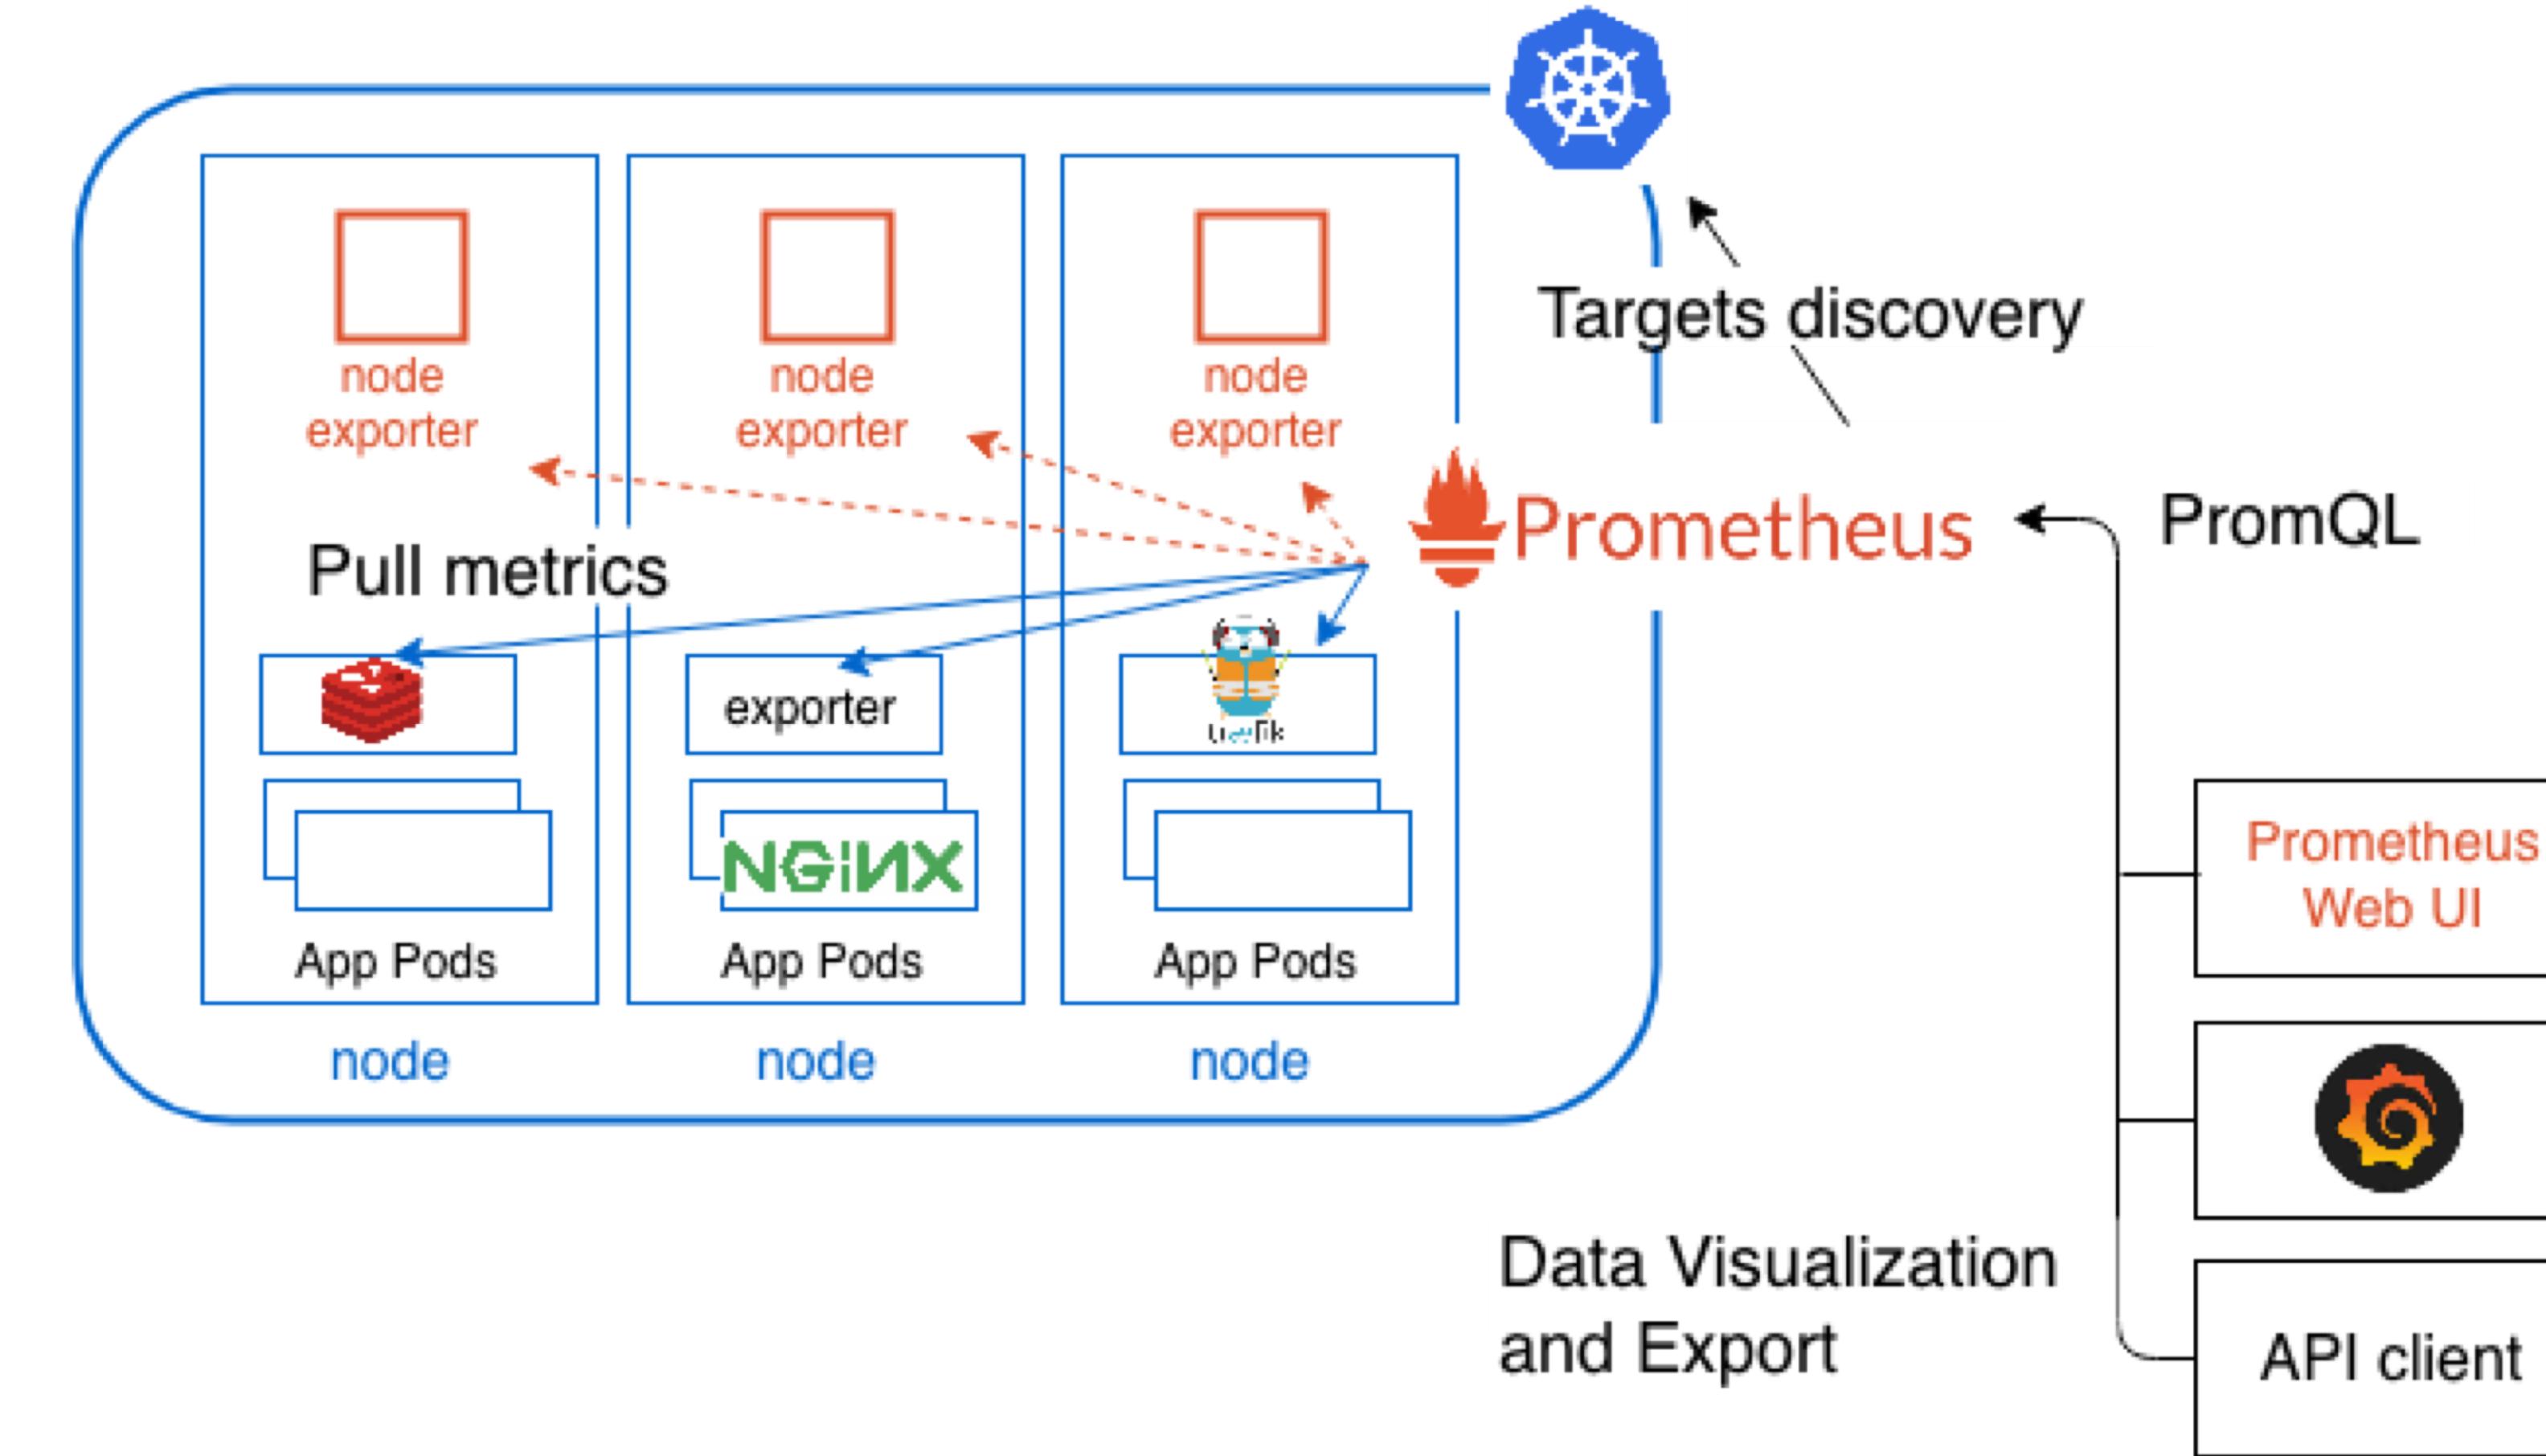
DEVOPS SPAIN

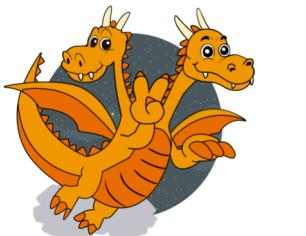
II EDICIÓN



2. Monitoring systems and applications

Metrics, scraping, querying,
graphing, prometheus-
native apps

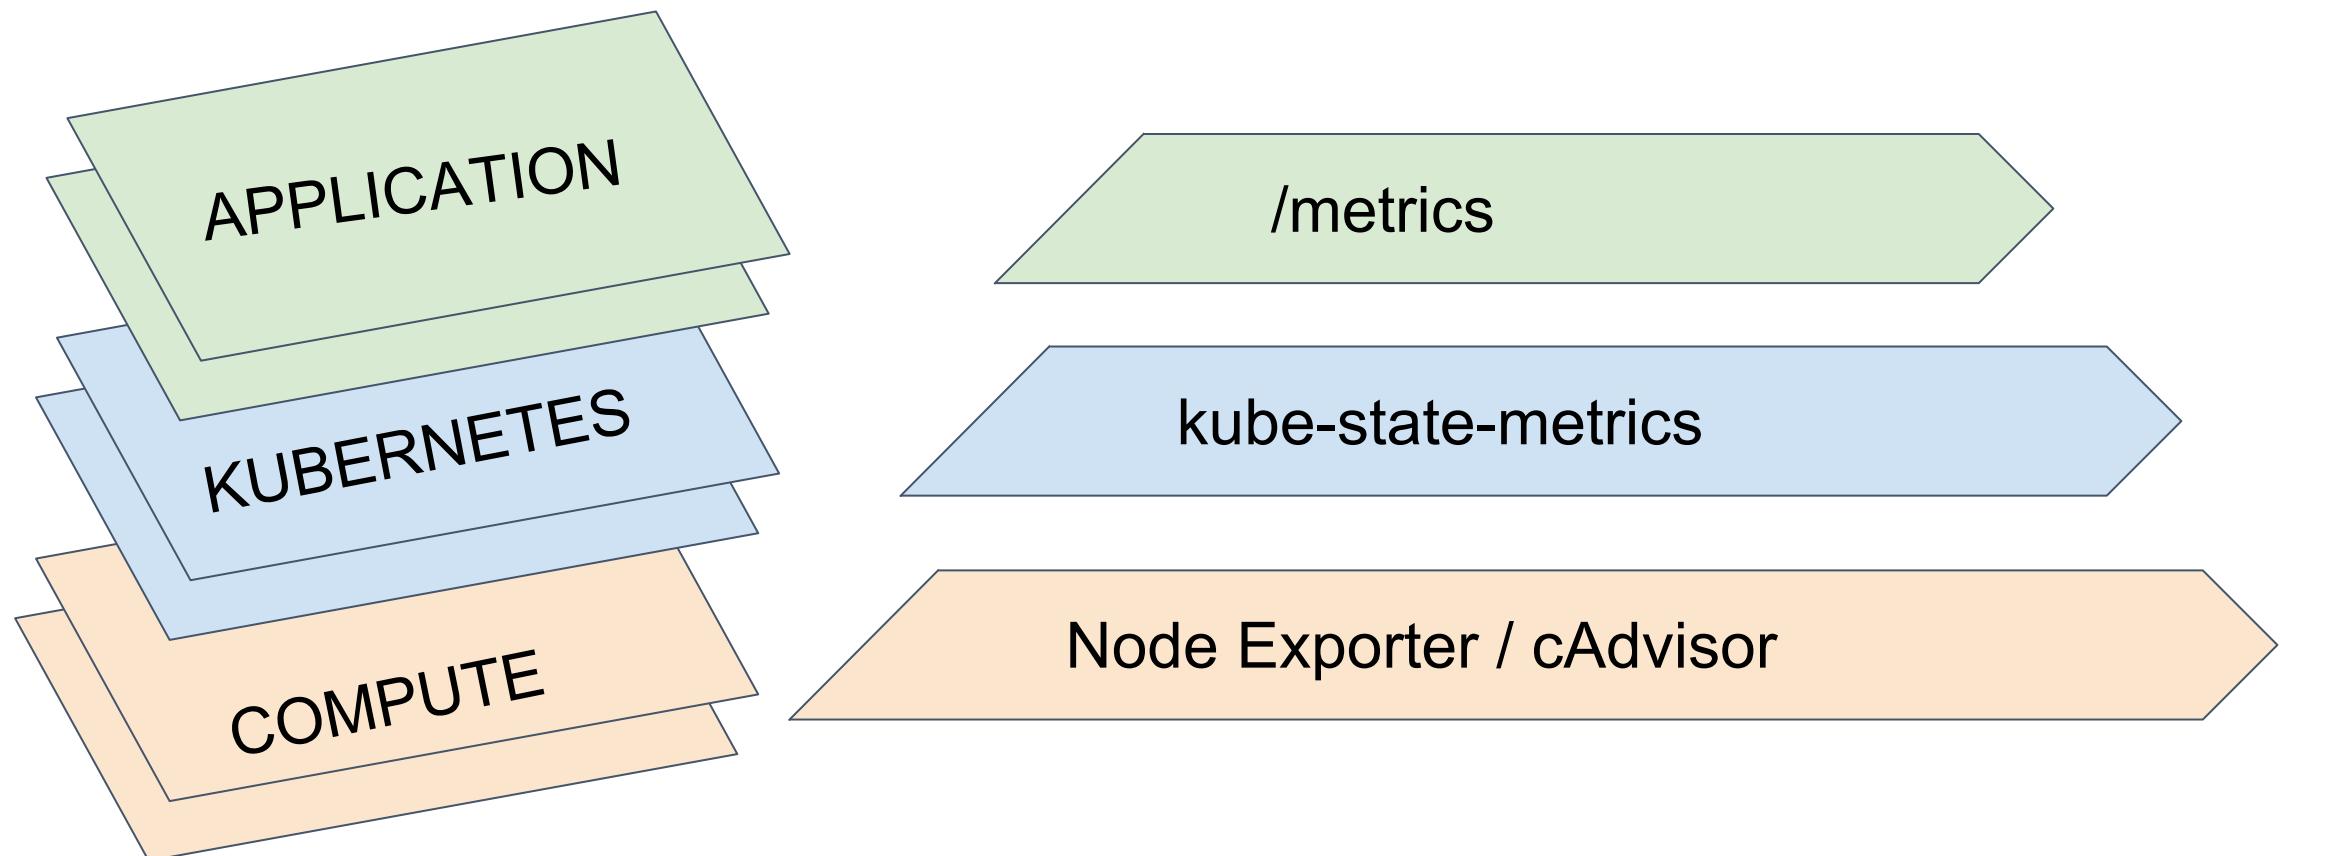




2.1 Monitoring systems

Available exporters

- node-exporter (Prometheus)
- kube-state-metrics (Kubernetes)
- cAdvisor (Google)
- Docker daemon (Docker)

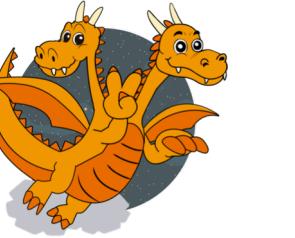


node-exporter metrics
arp, boottime, cpu, diskstats, entropy, exec, filesystem, ipvs, loadavg, meminfo, netclass, netstat, nfs, sockstat, stat, **textfile**, time, timex, vmstat, zfs ...

cAdvisor metrics
container_accelerator_memory_used_bytes, container_cpu_load_average_10s, container_fs_reads_total, container_last_seen, container_network_receive_errors_total

kube-state-metrics metrics
kube_cronjob_status_active, kube_namespace_annotations, kube_service_labels, kube_service_created, kube_replicaset_status_replicas, kube_node_status_condition .

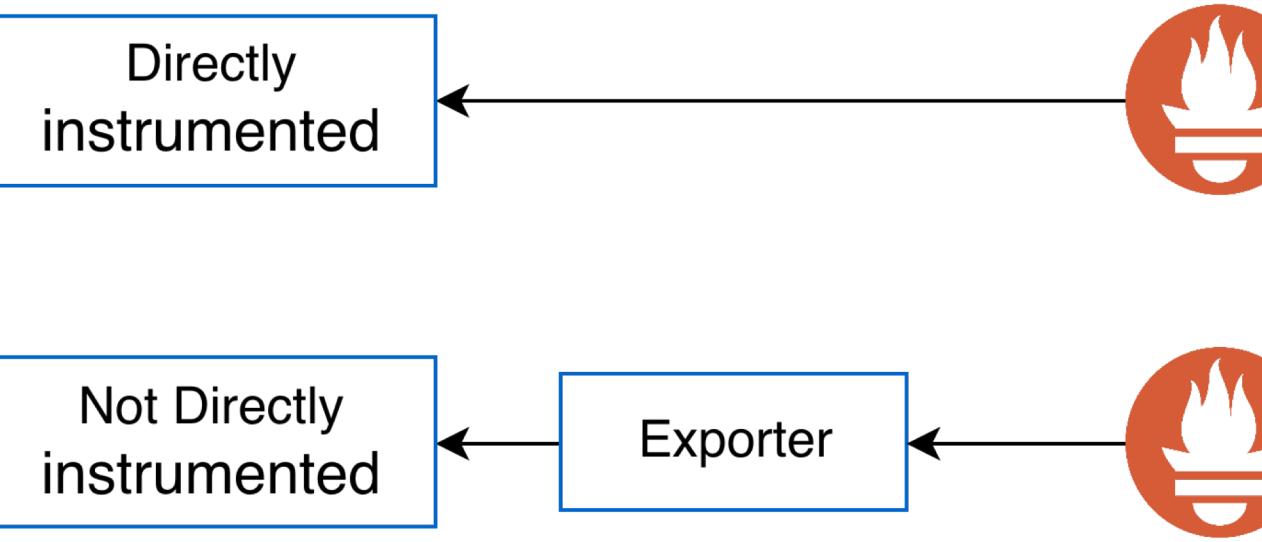
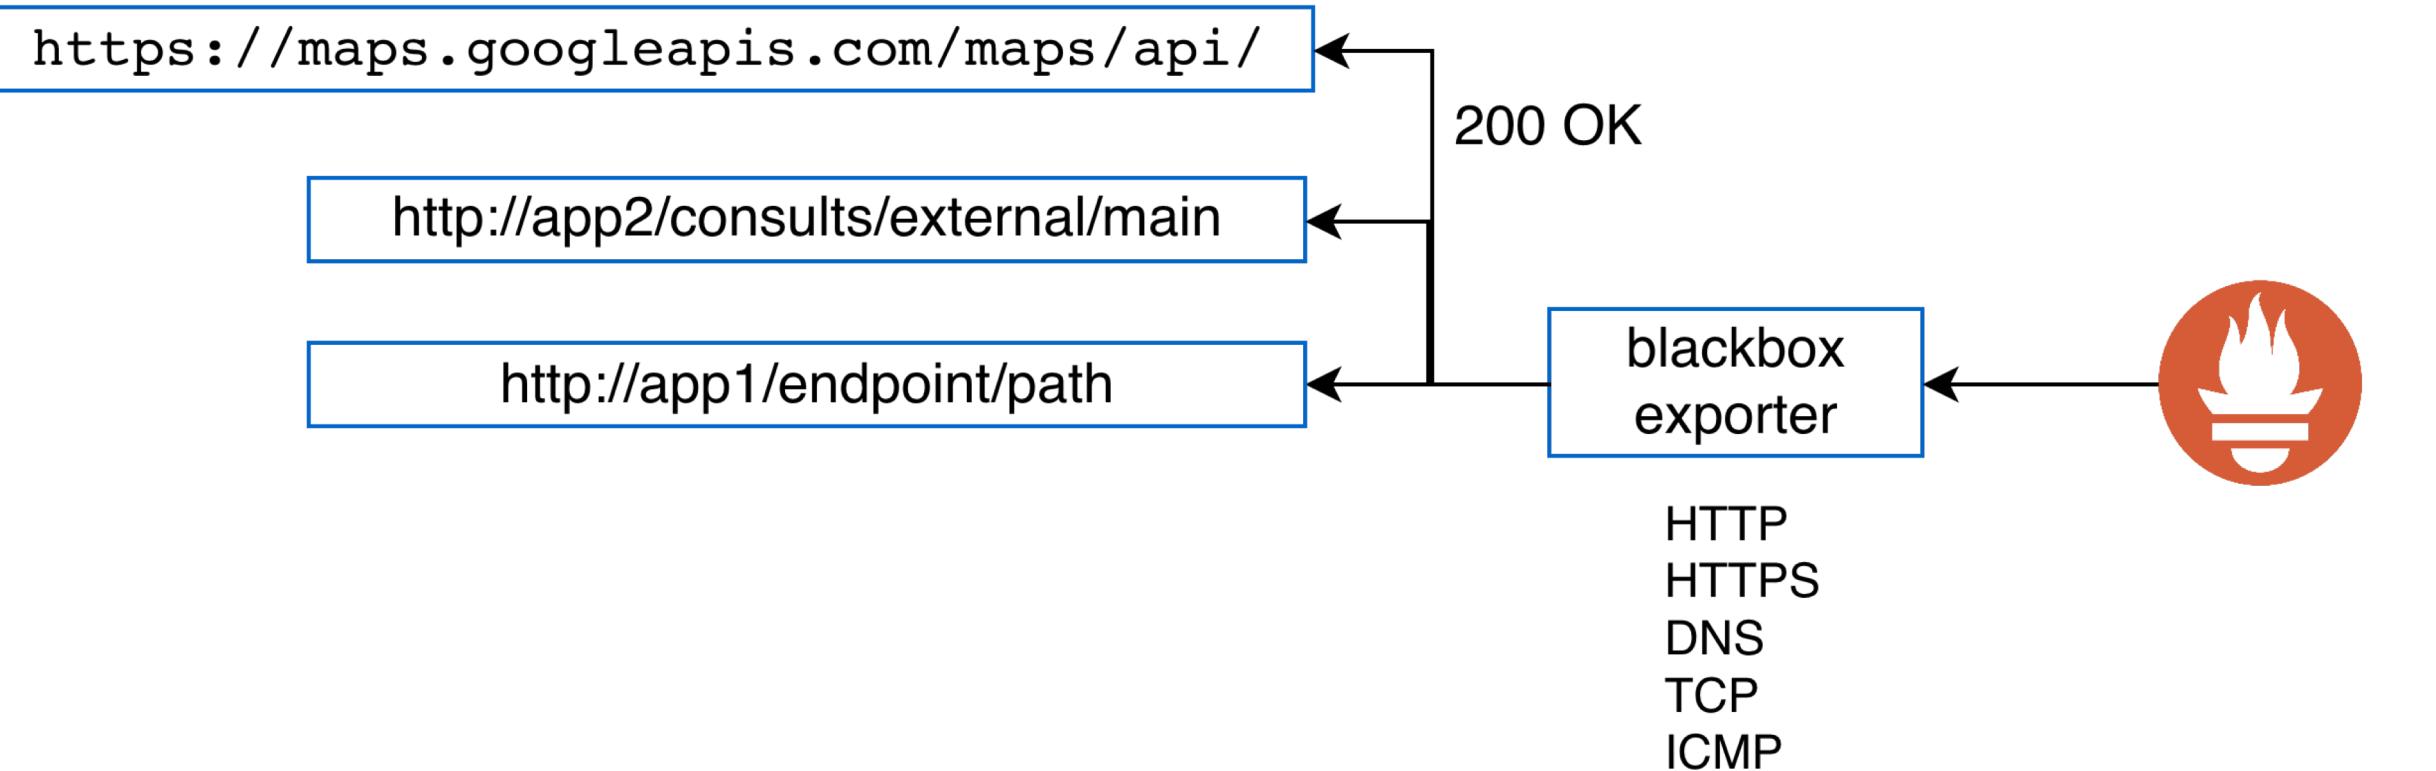
docker daemon metrics
engine_daemon_container_action_seconds_bucket, process_cpu_seconds_total

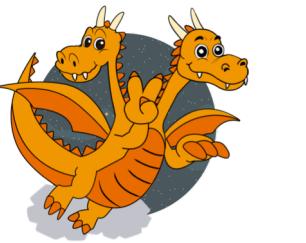


2.2 Monitoring applications

Application metrics

- Blackbox exporter: monitoring microservices endpoints
- Official exporters
- prometheus-native applications



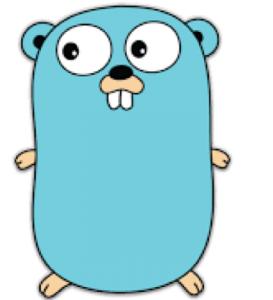


2.3 prometheus-native applications

Instrumenting applications

Before you can **monitor your services**, you need to add instrumentation to their code.

1. Implement the Prometheus metric types
2. Choose a Prometheus **client library** that matches the language of your application.
3. **Define internal metrics**
4. **Expose them via an HTTP endpoint** on your application instance.

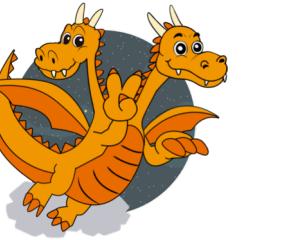


Available client libraries:

- Go
- Java or Scala
- Python
- Ruby
- Unofficial, third-party: bash, PHP, node.js,

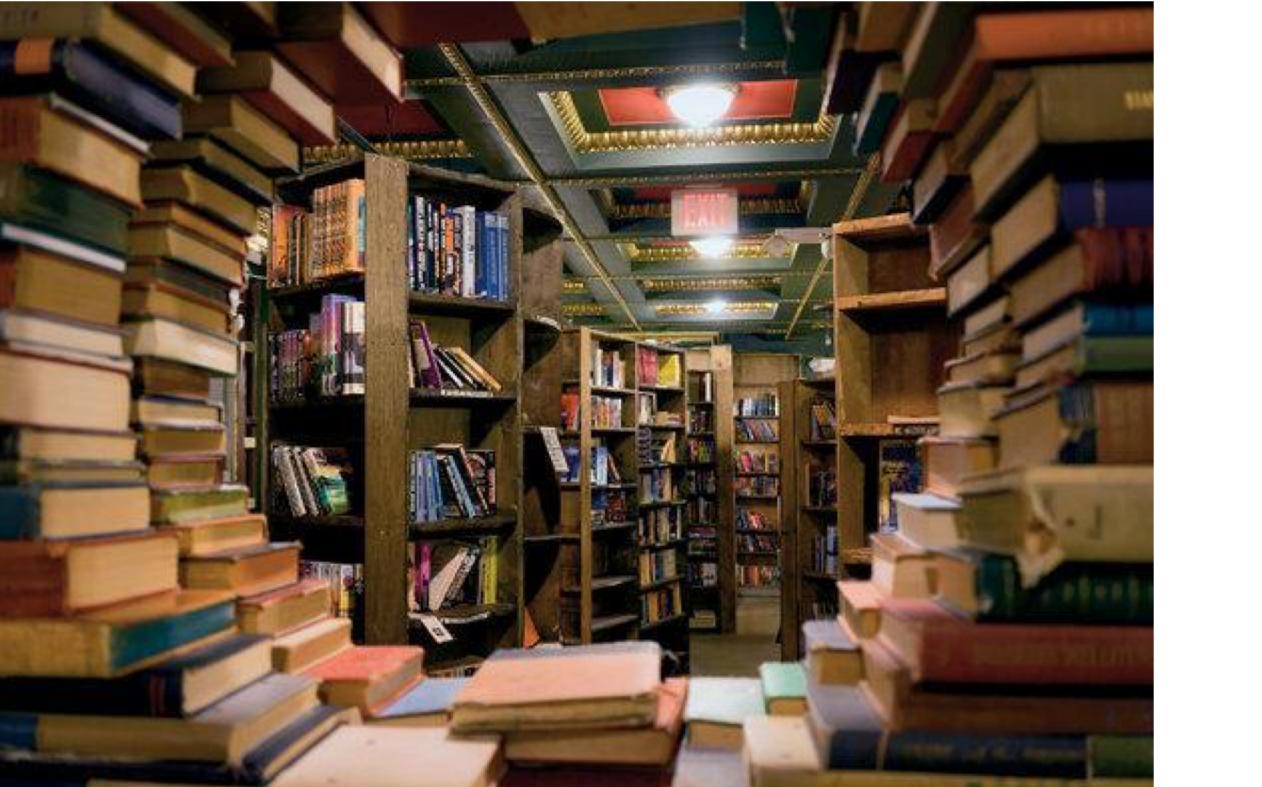
Available metrics types:

- Counter
- Gauge
- Histogram
- Summary



2.3.1 prometheus-native applications

Book store app



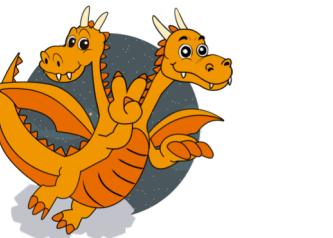
Possible actions taken by users:

- Buy a book

Every book categorized by:

- Genre: Terror, Romance, Science Fiction
- Pages

- I would like to know, **which genres** are buy the most by my clients.
- **Which editions** are preferred, pocket editions, portable..
- If a specific genre is most buy to read in a comfortable light edition.
- Number of **visits**
- **Time** spend before buying



2.3.1 prometheus-native applications

What metrics look like

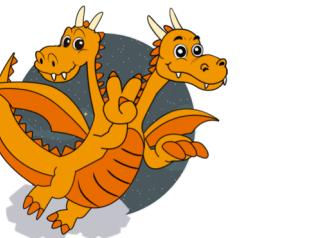
For every metric

```
# HELP books_sold Number of books sold
# TYPE books_sold counter
books_sold[genre="terror"] 199.0
books_sold[genre="romance"] 70.0
# HELP uptime uptime
# TYPE uptime gauge
uptime 4.2769899E7
# HELP systemload_average systemload_average
# TYPE systemload_average gauge
systemload_average 0.55
# HELP heap_committed heap_committed
# TYPE heap_committed gauge
heap_committed 2234880.0
```

Description

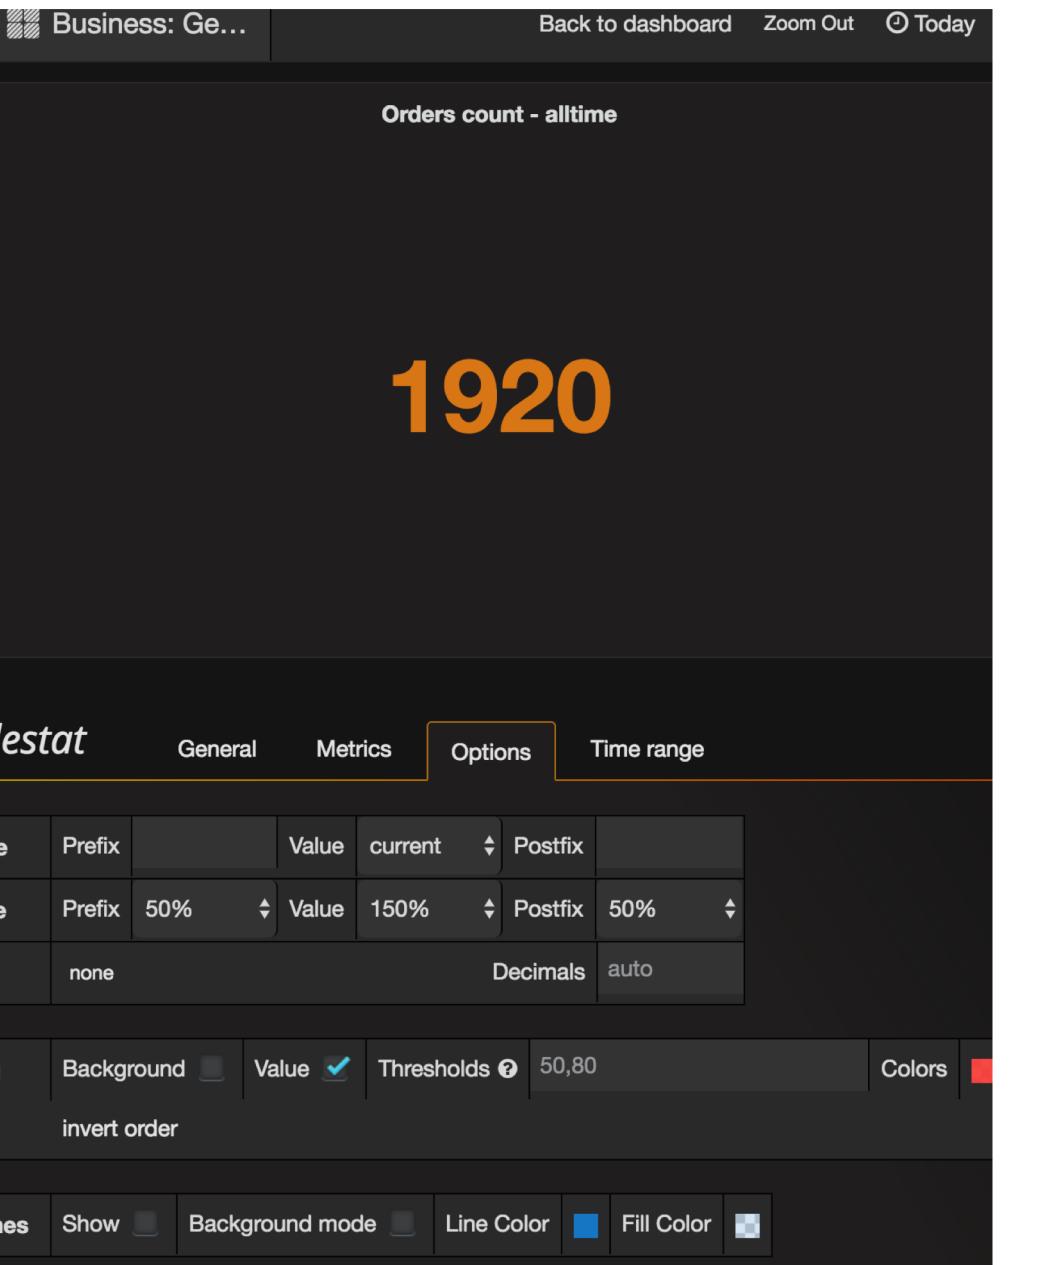
Metric type

Values are float64



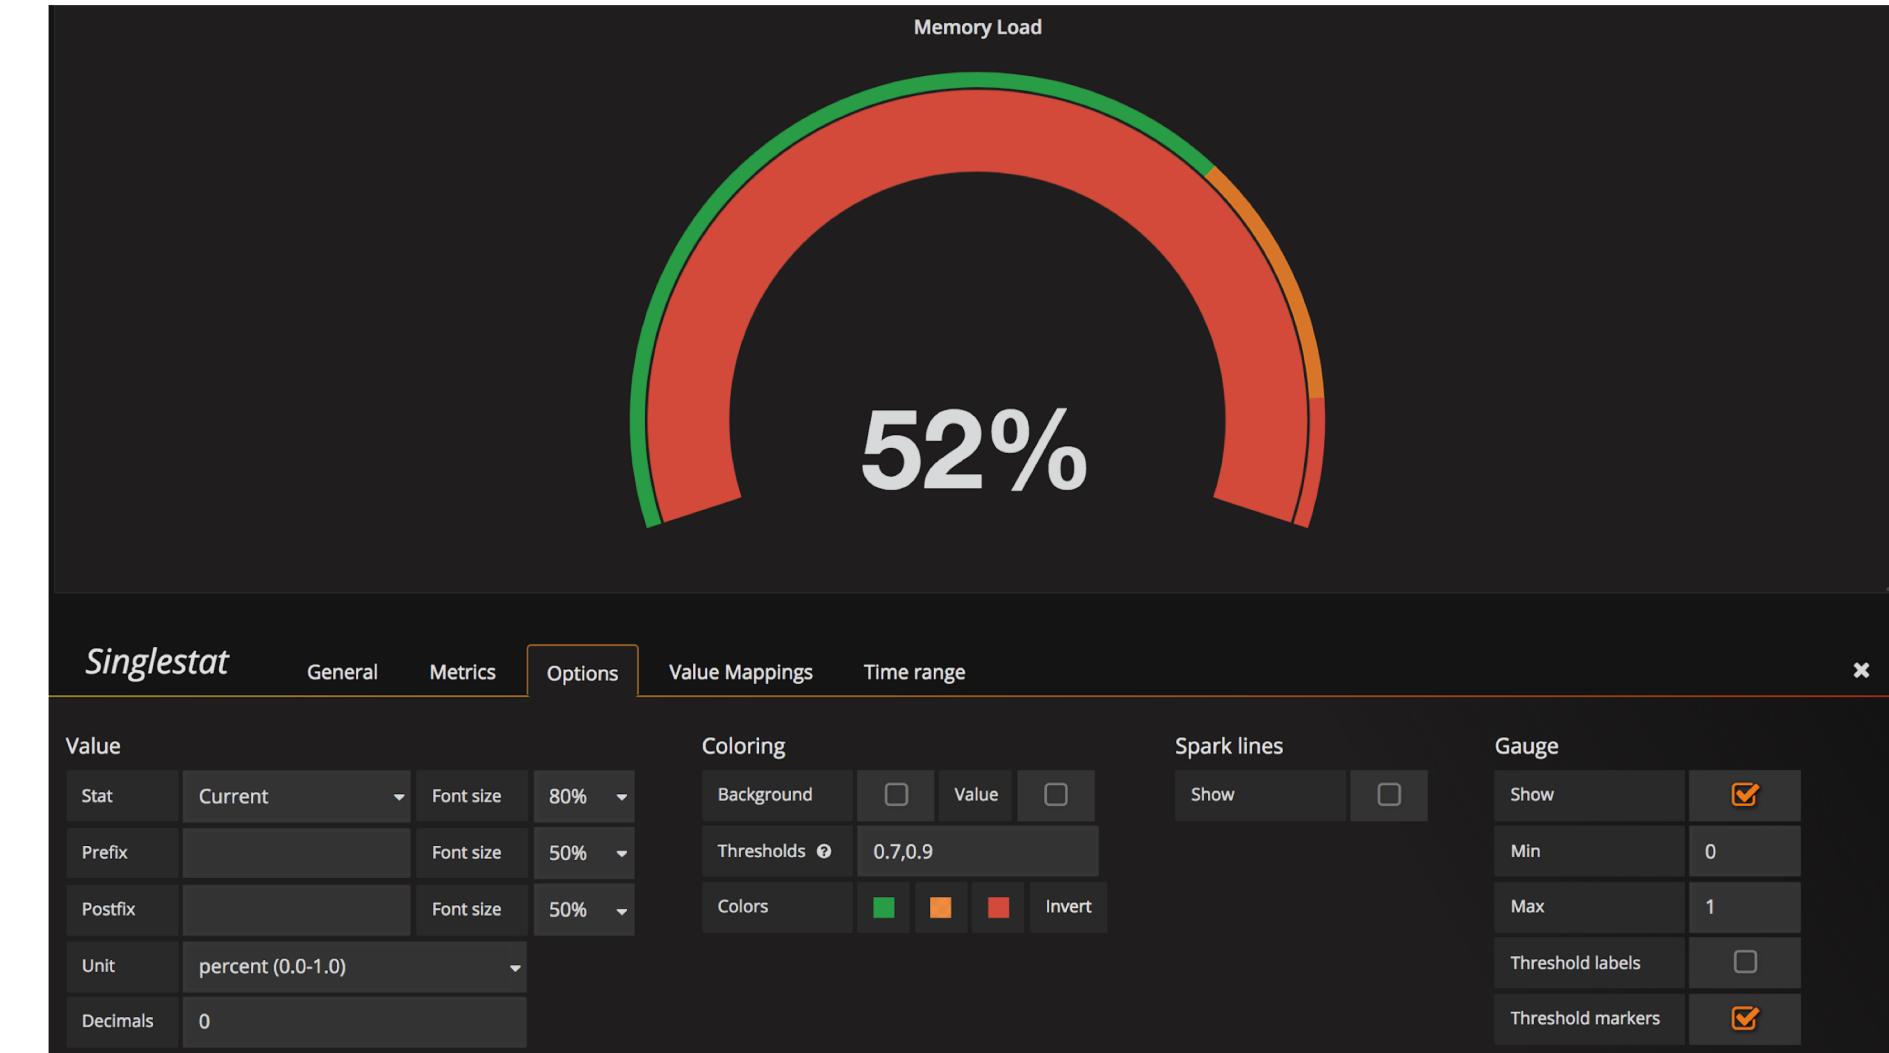
2.3.2 prometheus-native applications

Metrics types: Counter and Gauge



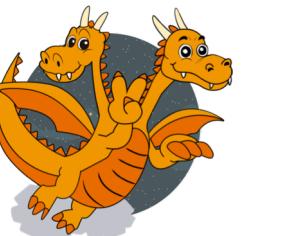
Counter:

- Increment or reset
- Number of visits or books sold



Gauge:

- Type of counter
- Increment or decrement
- Common use: represent memory or CPU



2.3.2 prometheus-native applications

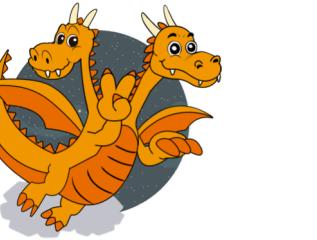
Metrics types: Counter

```
# HELP books_sold Number of books sold
# TYPE books_sold counter
books_sold{genre="terror"} 2
books_sold{genre="romance"} 1
```



```
booksSold = promauto.NewCounterVec(prometheus.CounterOpts{
    Name: "books_sold",
    Help: "The total number of books sold",
},
[]string{"genre"})
```

```
r.POST("/buy/:genre", func(c *gin.Context) {
    genre := c.Param("genre")
    booksSold.WithLabelValues(genre).Inc()
    c.JSON(200, gin.H{
        "message": "Thank you for purchasing a book",
    })
})
```



2.3.2 prometheus-native applications

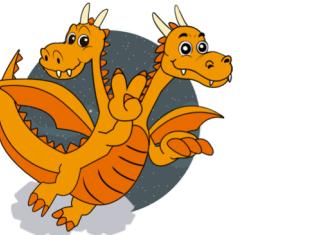
Metrics types: Gauge

```
# HELP online_users Users that are online
# TYPE online_users gauge
online_users 3
```

```
onlineUsers = promauto.NewGauge(prometheus.GaugeOpts{
    Name: "online_users",
    Help: "Users that are online",
})
```

```
r.GET("/", func(c *gin.Context) {
    onlineUsers.Inc()
    c.JSON(200, gin.H{
        "message": "Look at all the books we have!",
    })
})
```

```
r.POST("/buy/:genre", func(c *gin.Context) {
    onlineUsers.Dec()
    c.JSON(200, gin.H{
        "message": "Thank you for purchasing a book",
    })
})
```

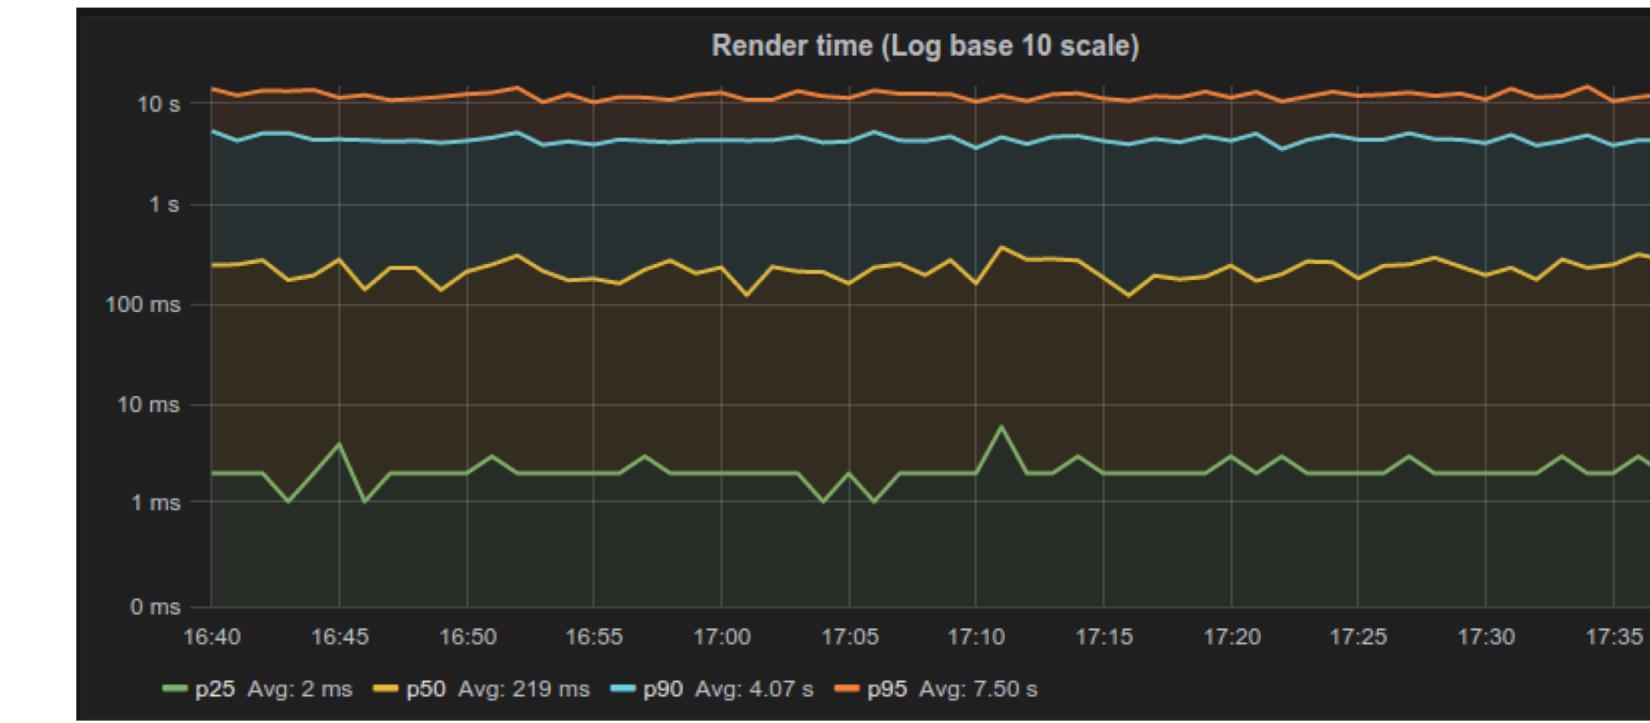
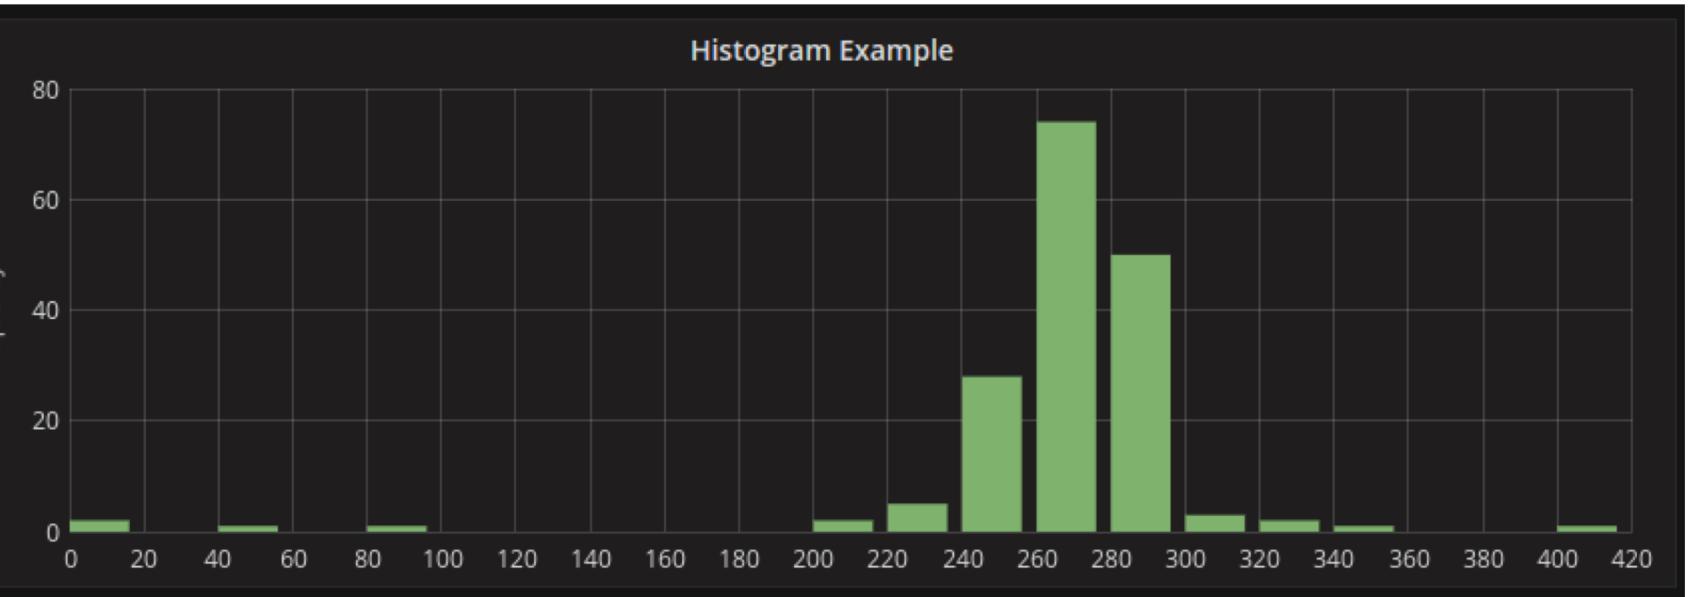


2.3.3 prometheus-native applications

Metrics types: Histogram and Summary

Histogram:

- Values grouped in buckets
- **Cumulative**
- Number of books with number of pages ranged from 50 to 100 pages.



Summary:

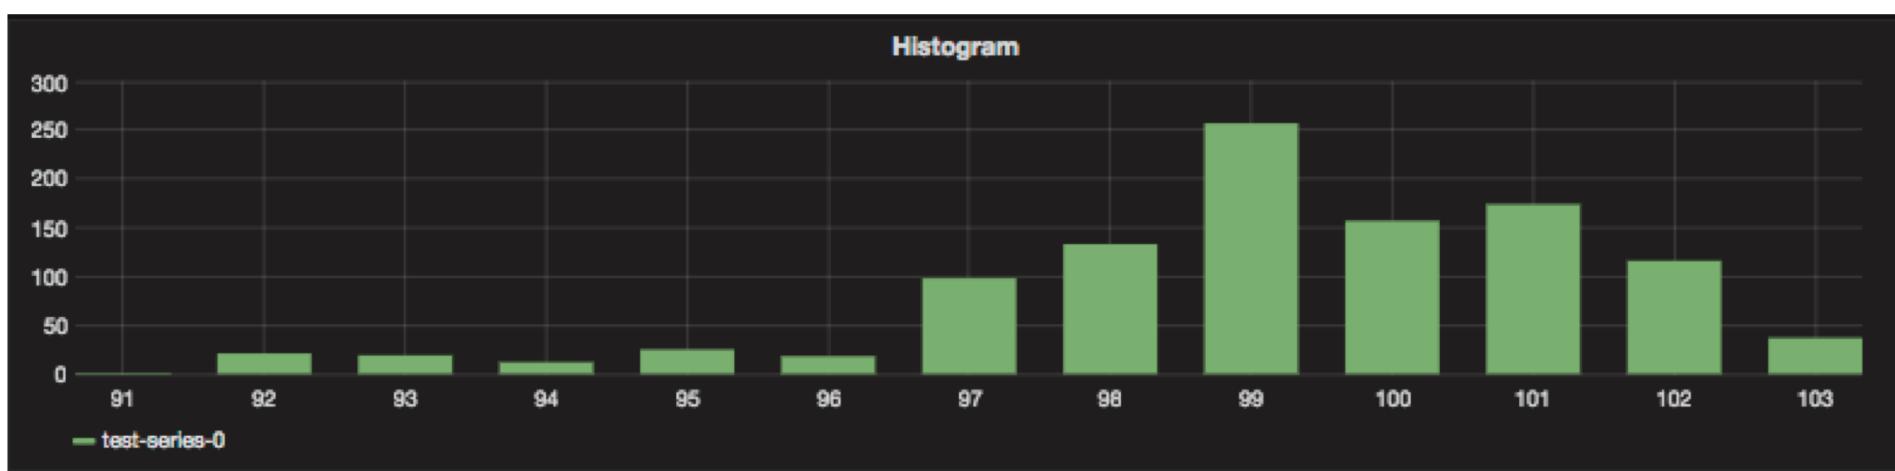
- **Average and Percentiles** information
- Get extremes behaviours
- Time spent in the application previous to buy anything

2.3.3 prometheus-native applications

Metrics types: Histogram

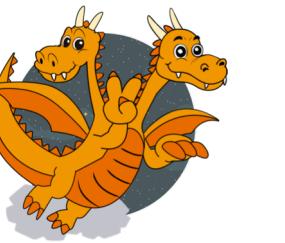
```
# HELP books_sold_pages Books sold with number of pages
# TYPE books_sold_pages histogram
books_sold_pages_bucket{genre="terror",le="20"} 5
books_sold_pages_bucket{genre="terror",le="100"} 21
books_sold_pages_bucket{genre="terror",le="150"} 29
books_sold_pages_bucket{genre="terror",le="200"} 32
books_sold_pages_bucket{genre="terror",le="+Inf"} 32
books_sold_pages_sum{genre="terror"} 2624
books_sold_pages_count{genre="romance"} 32
```

```
bSoldPages = promauto.NewHistogramVec(
    prometheus.HistogramOpts{
        Name:      "books_sold_pages",
        Help:      "Books sold with number of pages",
        Buckets:  []float64{20.0, 100.0, 150.0, 200.0},
    },
    []string{"genre"})
```



```
r.POST("/buy/:genre", func(c *gin.Context) {
    genre := c.Param("genre")
    bSoldPages.WithLabelValues(genre).Observe(rand.Float64() * 200)

    c.JSON(200, gin.H{
        "message": "Thank you for purchasing a book",
    })
})
```



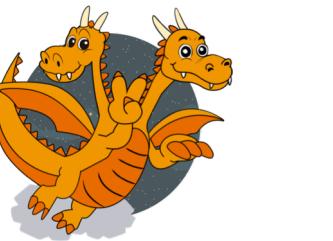
2.3.3 prometheus-native applications

Metrics types: Summary

```
# HELP time_spent Time spent to buy a book
# TYPE time_spent summary
time_spent{genre="romance",quantile="0.5"} 36.08
time_spent{genre="romance",quantile="0.9"} 68.68
time_spent{genre="romance",quantile="0.99"} 94.05
time_spent_sum{genre="romance"} 912.74
time_spent_count{genre="romance"} 22
```

```
timeSpent = promauto.NewSummaryVec(prometheus.SummaryOpts{
    Name:      "time_spent",
    Help:      "Time spent to buy a book",
    Objectives: map[float64]float64{0.5: 0.05, 0.9: 0.01, 0.99:
        0.001},
    },
    []string{"genre"})
```

```
r.POST("/buy/:genre", func(c *gin.Context) {
    genre := c.Param("genre")
    timeSpent.WithLabelValues(genre).Observe(rand.Float64()*100)
    c.JSON(200, gin.H{
        "message": "Thank you for purchasing a book",
    })
})
```



2.3.4 prometheus-native applications

Book Store insights



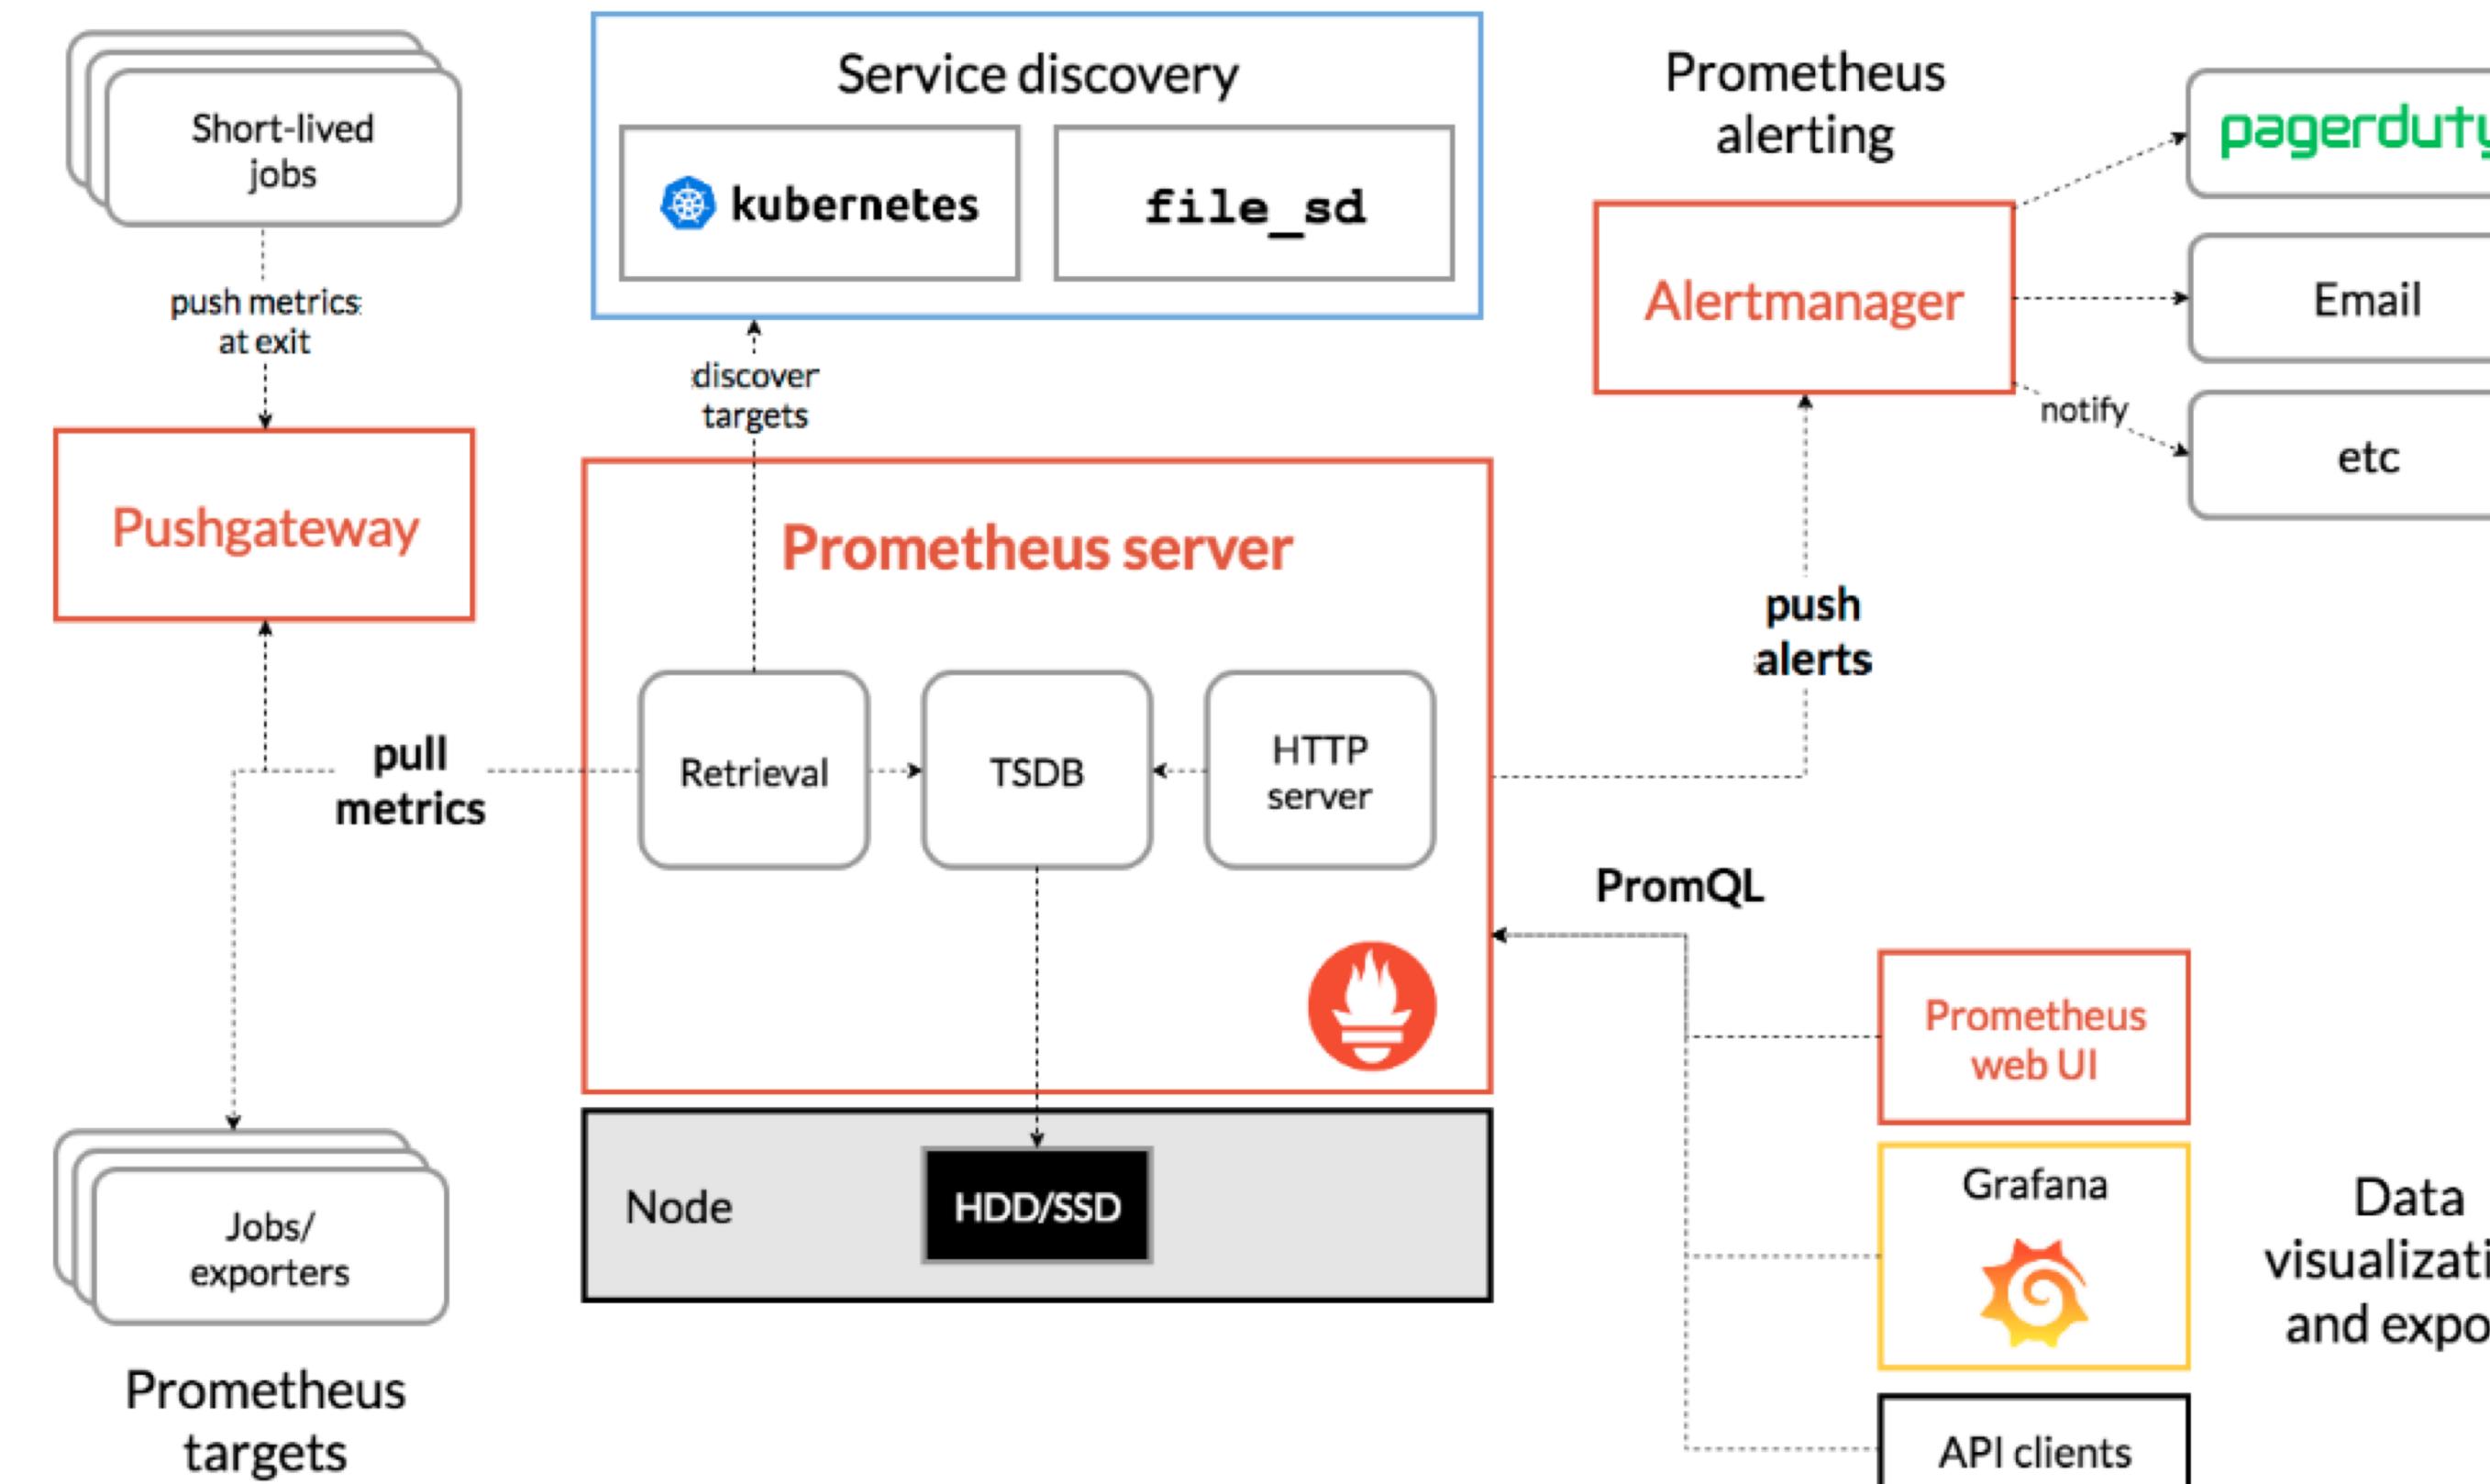
DEVOPS SPAIN

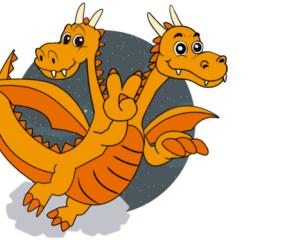
II EDICIÓN



3. Alerting and Integrations

Grouping, silencing, alerting

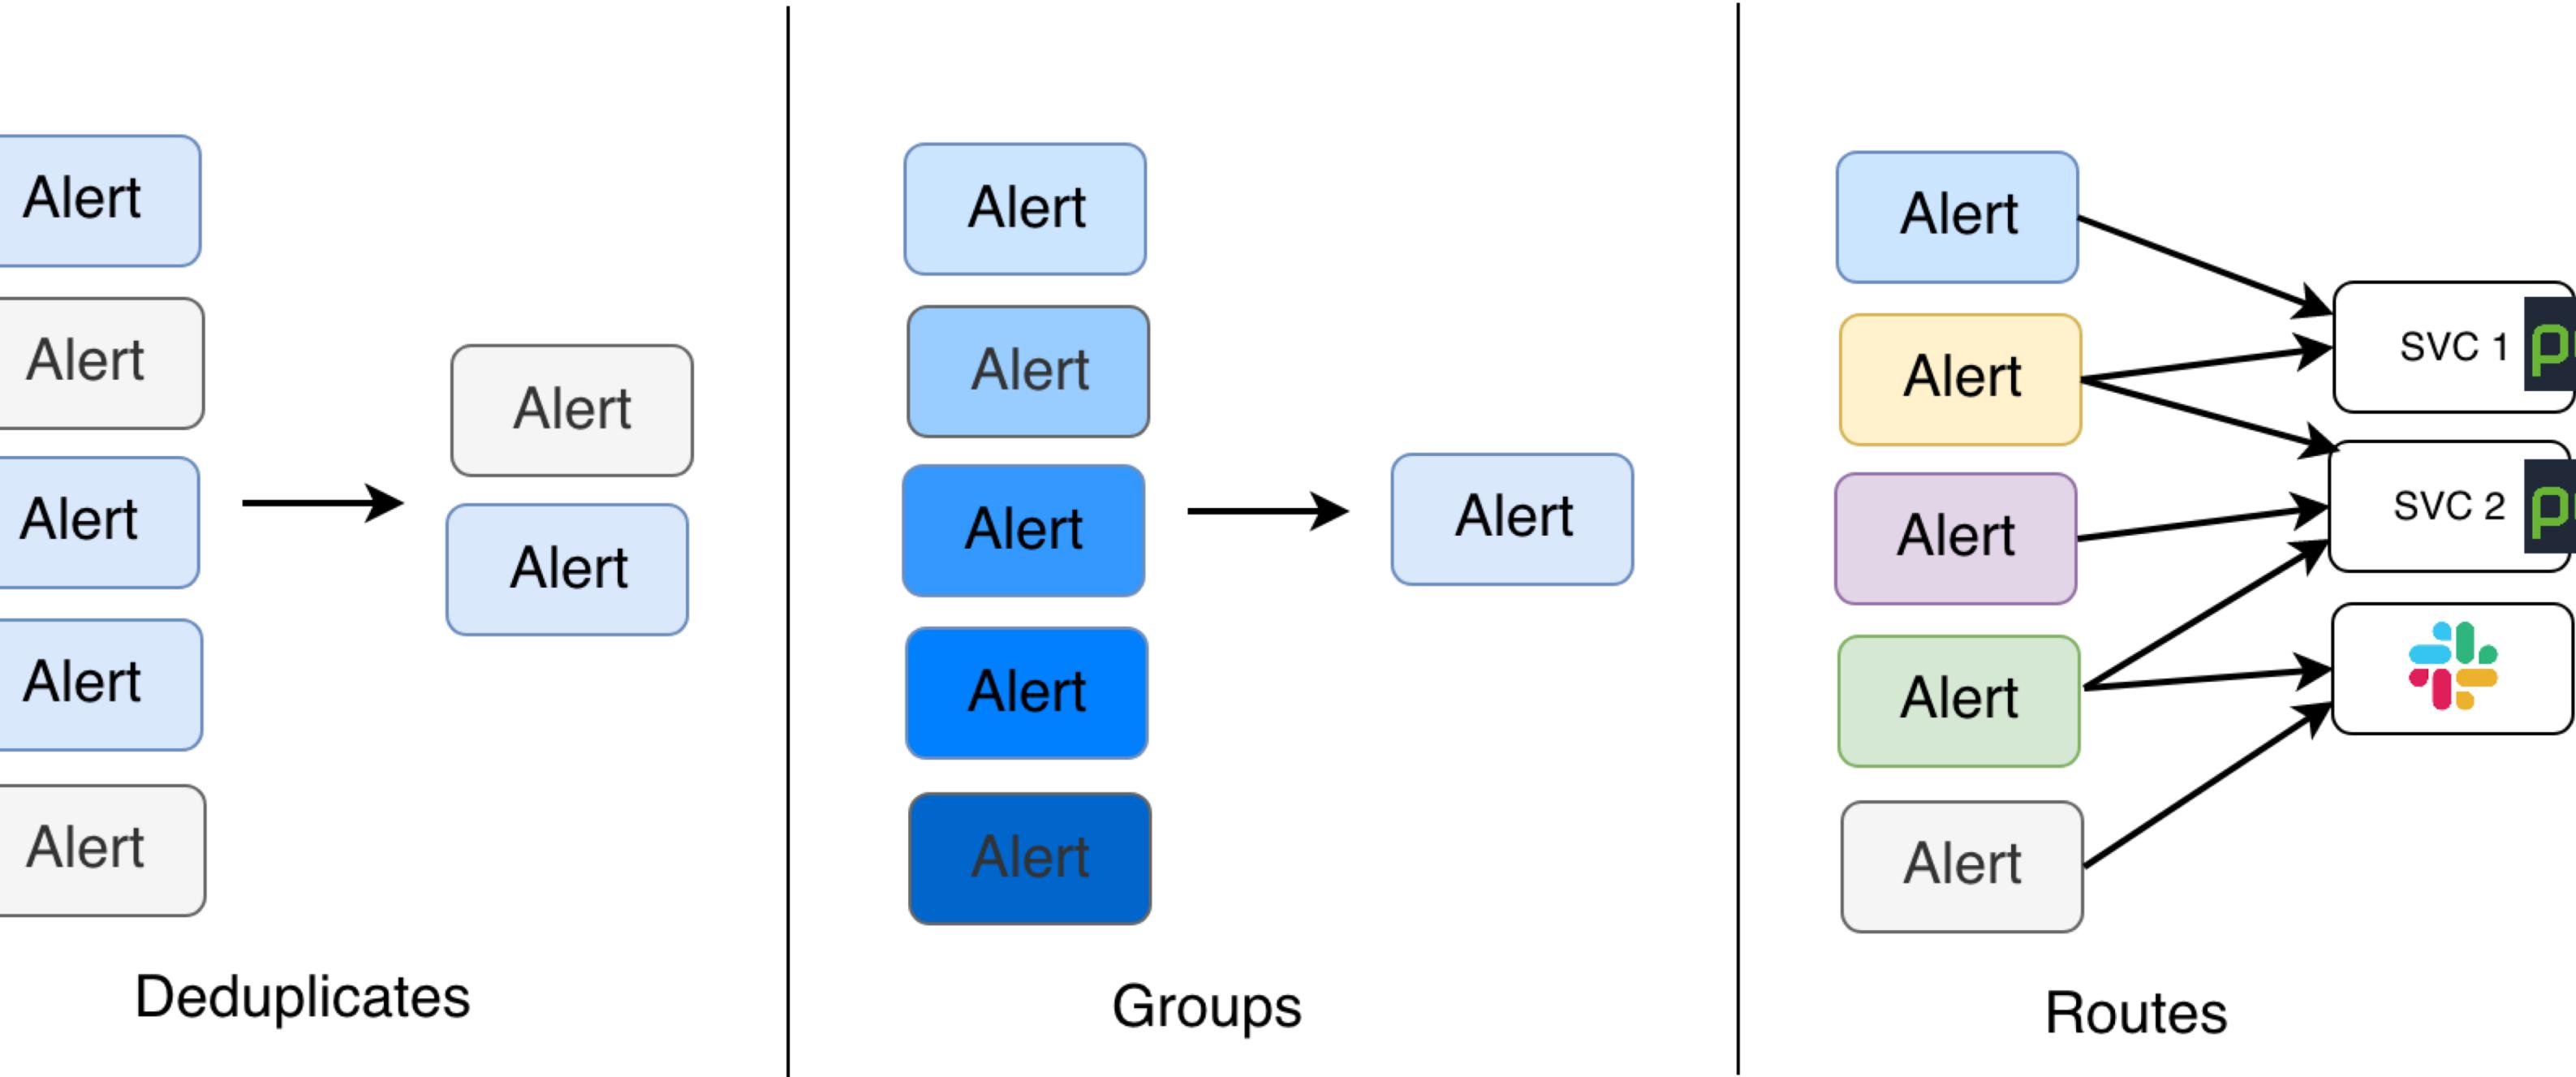




3.1 Alertmanager

Alerts notification management

- Deduplicates
- Groups
- Routes
- Inhibits
- Silence



DEVOPS SPAIN

II EDICIÓN

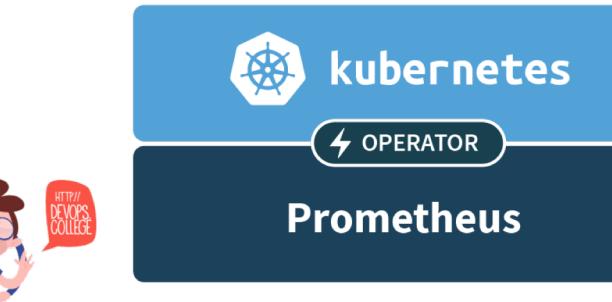


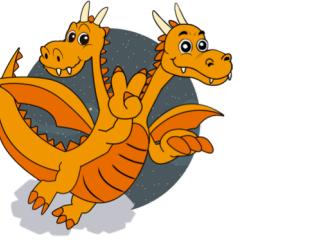
4. Scaling and Ecosystem

Federation, Thanos, K8s operator, Signifai

Getting the most out of Prometheus in real life case scenarios.

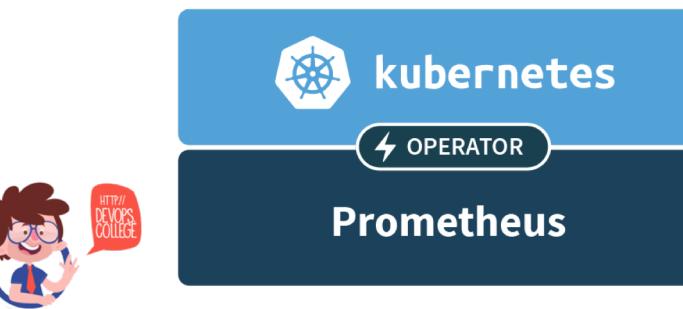
- New necessities are born: **Scaling** Prometheus
- A whole ecosystem appears to:
 - **Enable** an easy deployment, management and operation (Helm, Operator, ...)
 - **Complement** functionalities and support new approaches to cover specific necessities (Thanos, M3,Cortex ..)
 - Provide new paths to **exploit** the observability and instrumentations of our applications (Signifai)





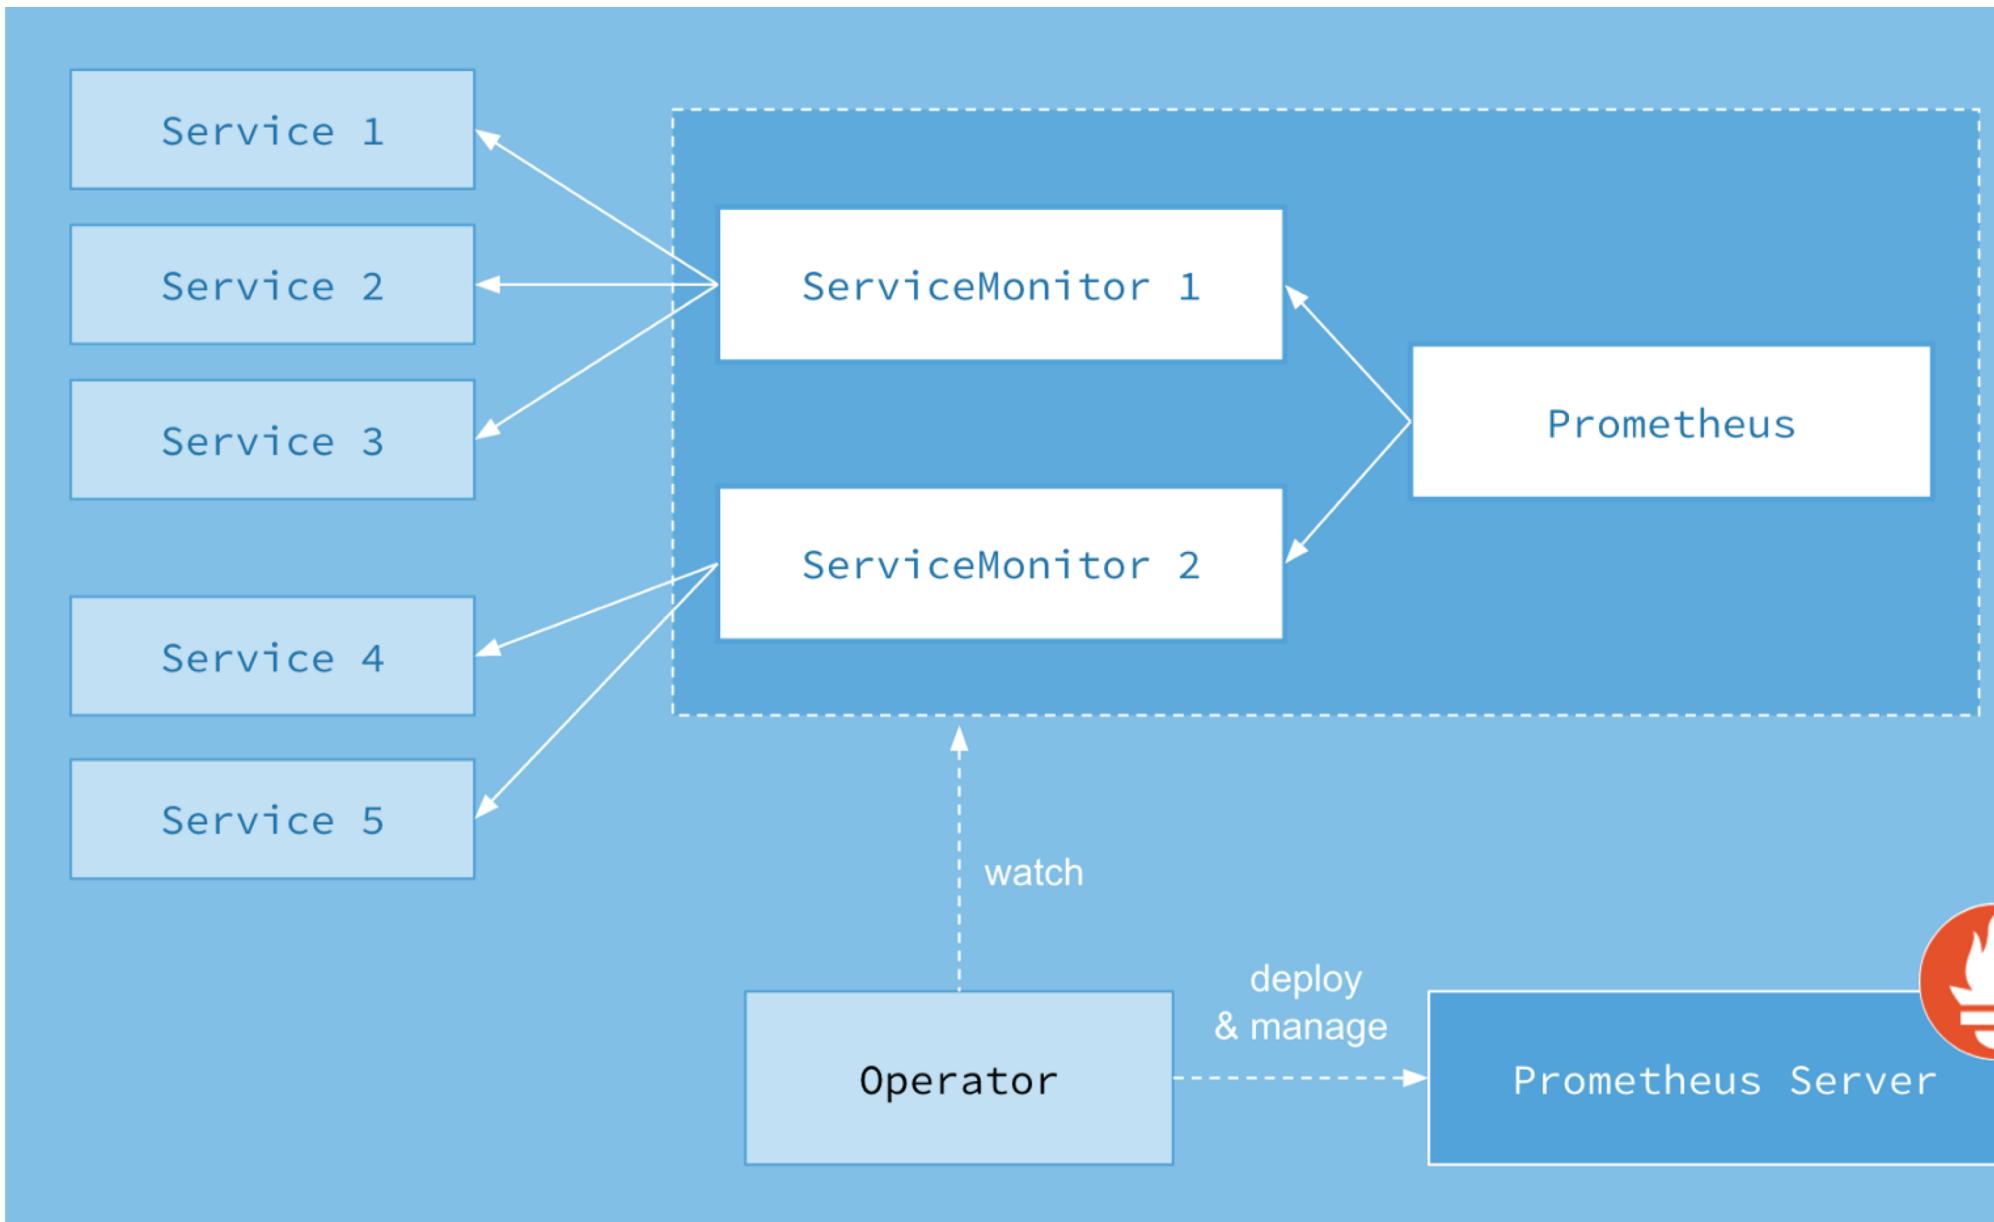
4.1 Deploying Prometheus

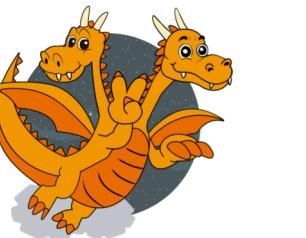
K8s prometheus-operator



Custom resources:

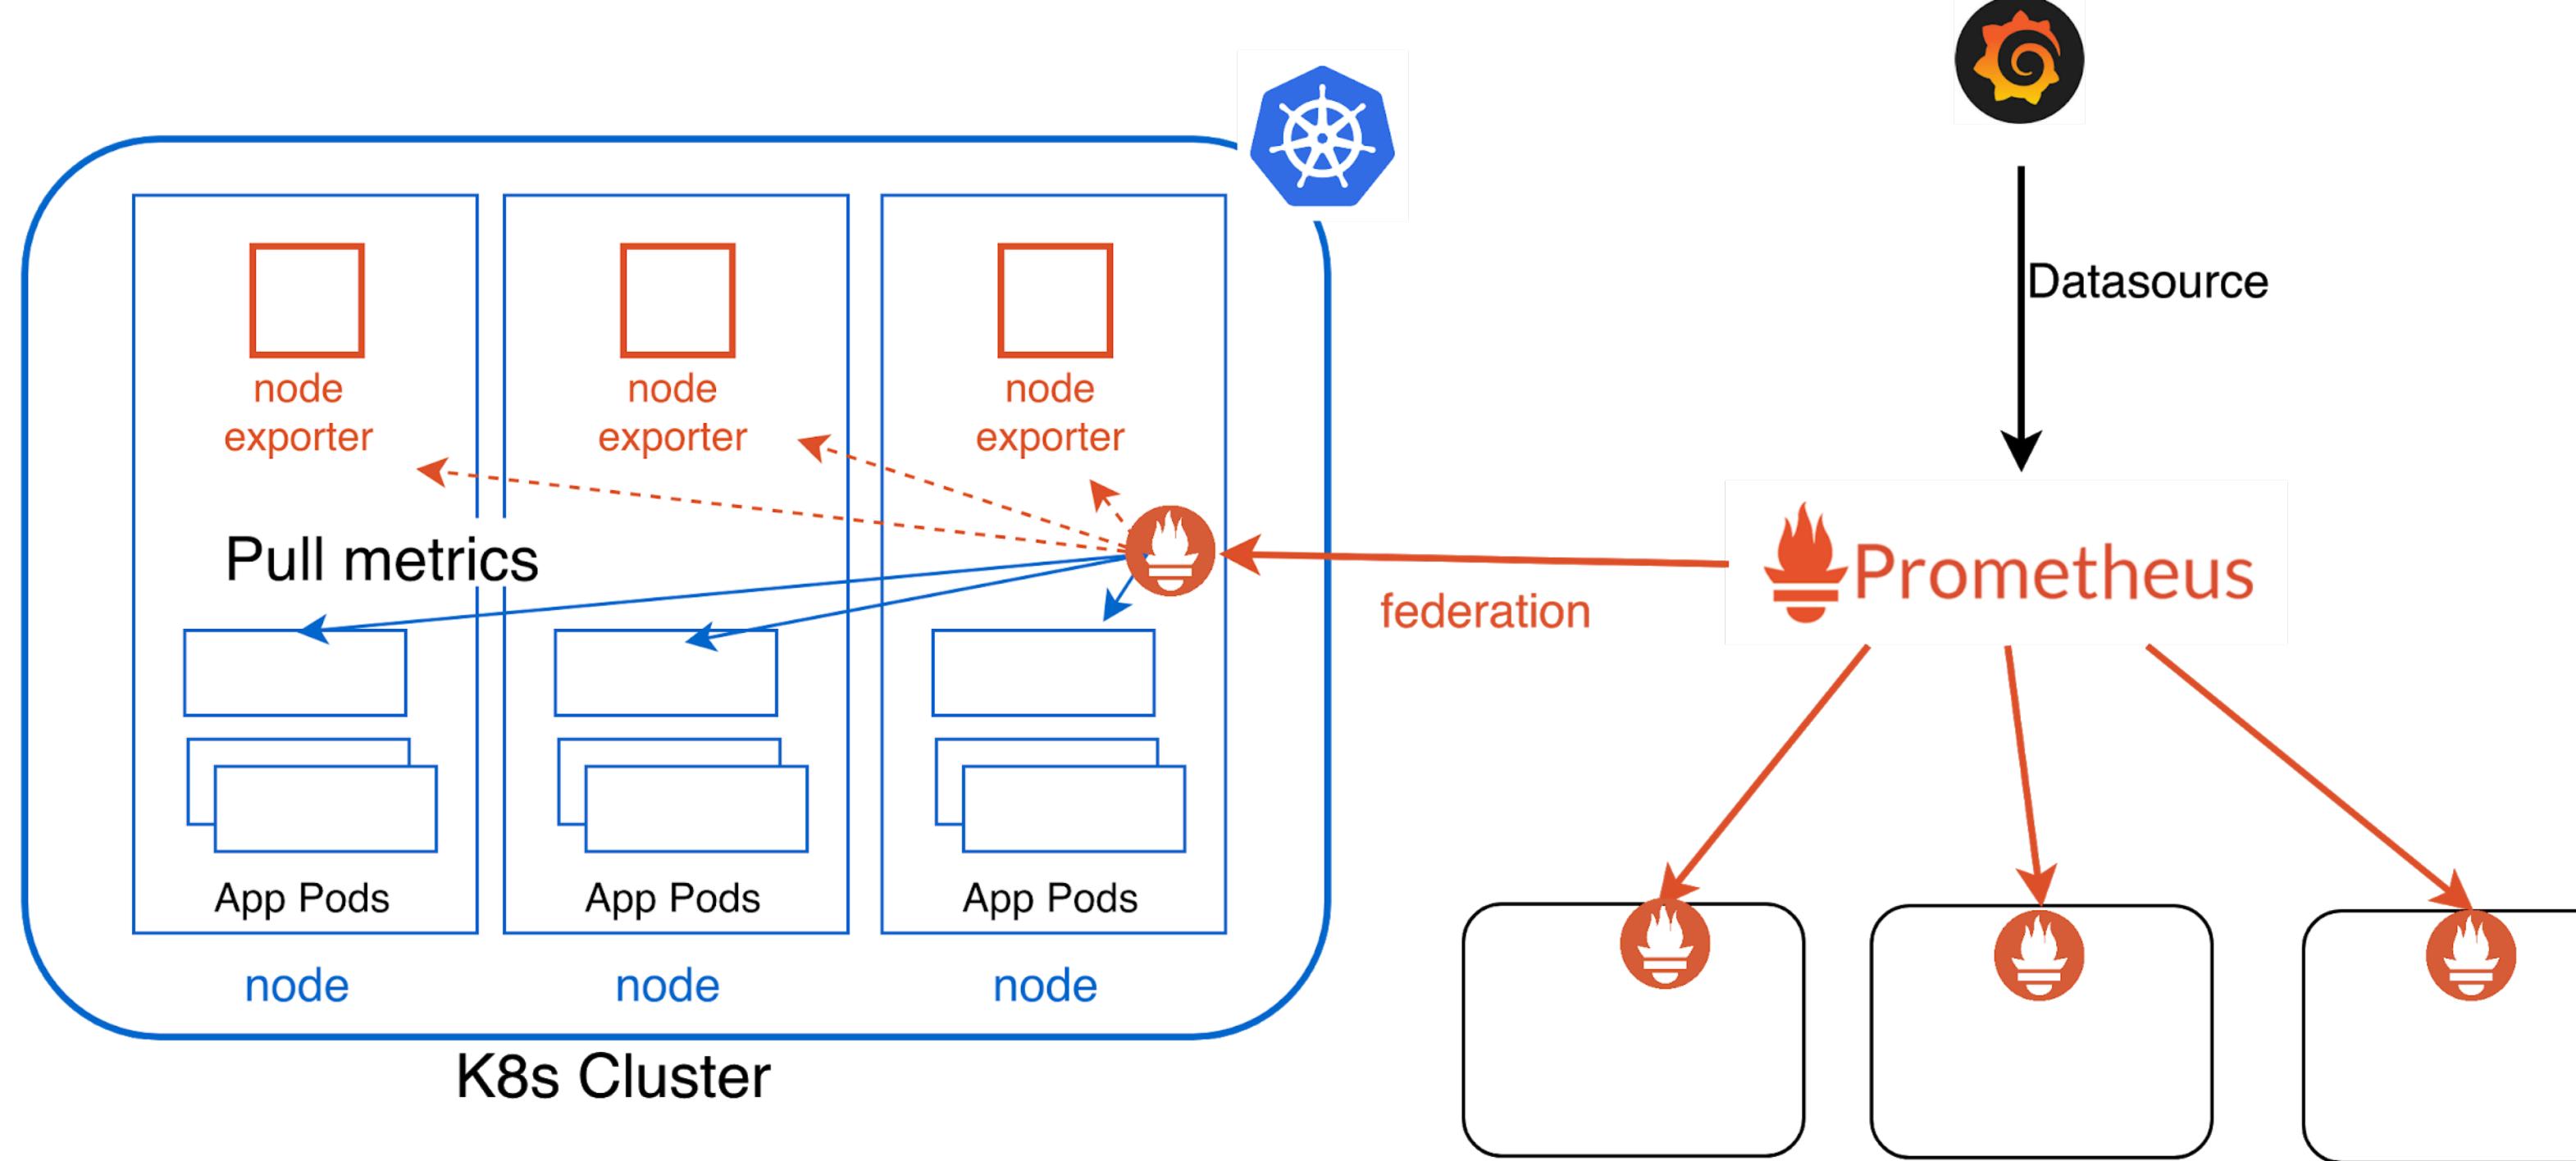
- **Prometheus**
- **ServiceMonitor**
- **Alertmanager**
- **PrometheusRule**

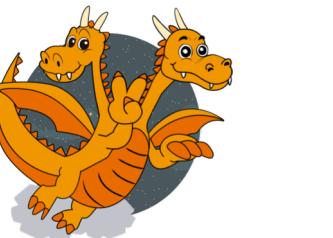




4.2 Scaling Prometheus

Federation

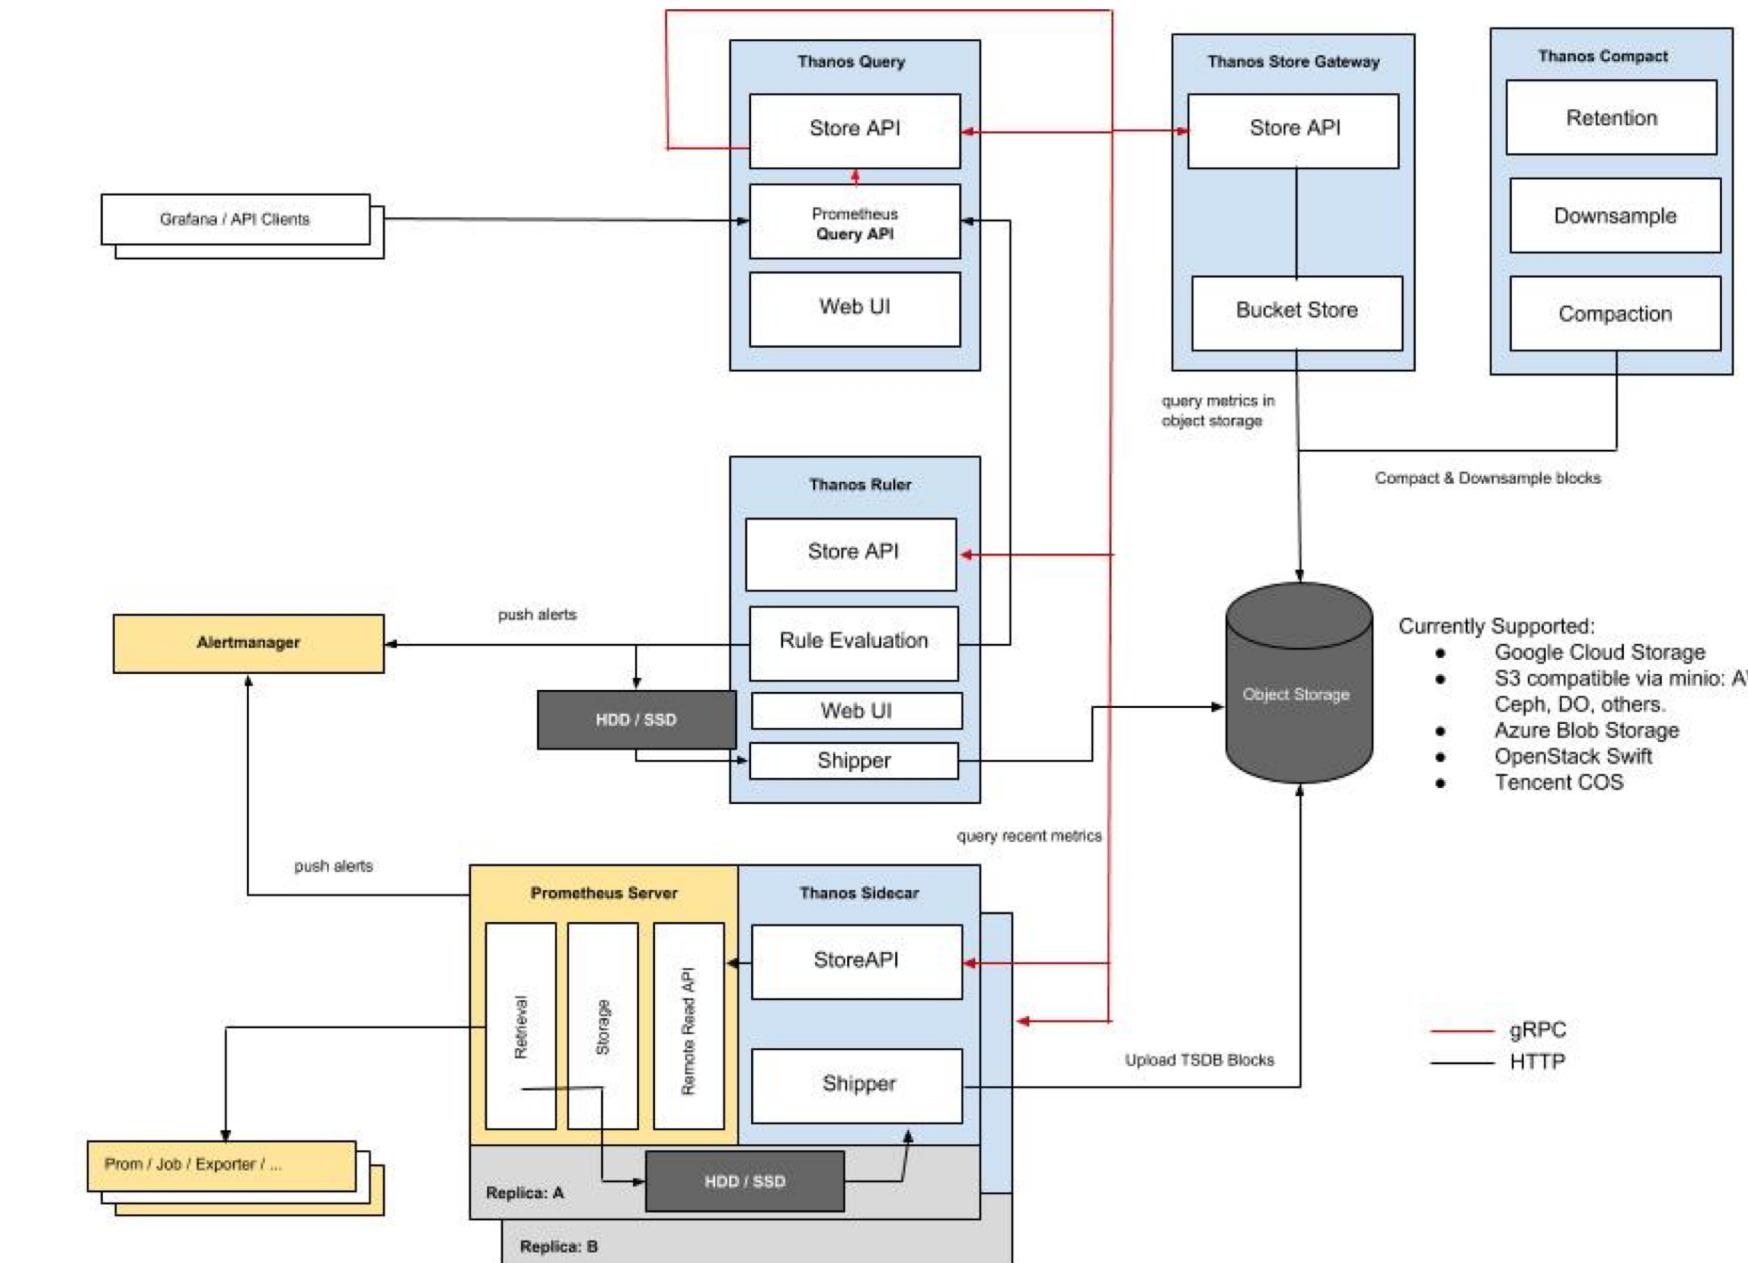
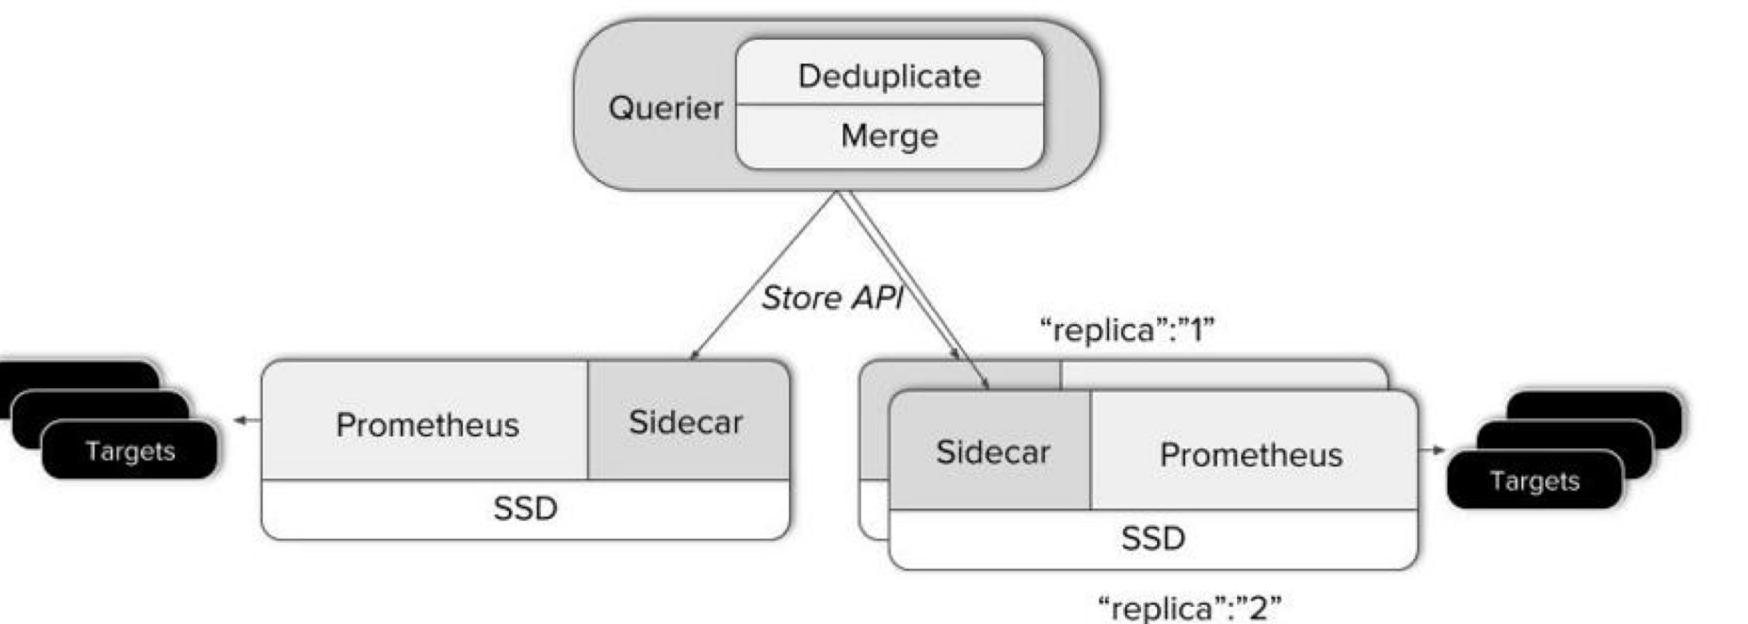


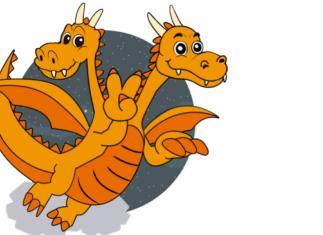


4.3 Thanos

Highly available Prometheus setup with long term storage capabilities

- Global view
- **Distributed queries**
- Unlimited retention

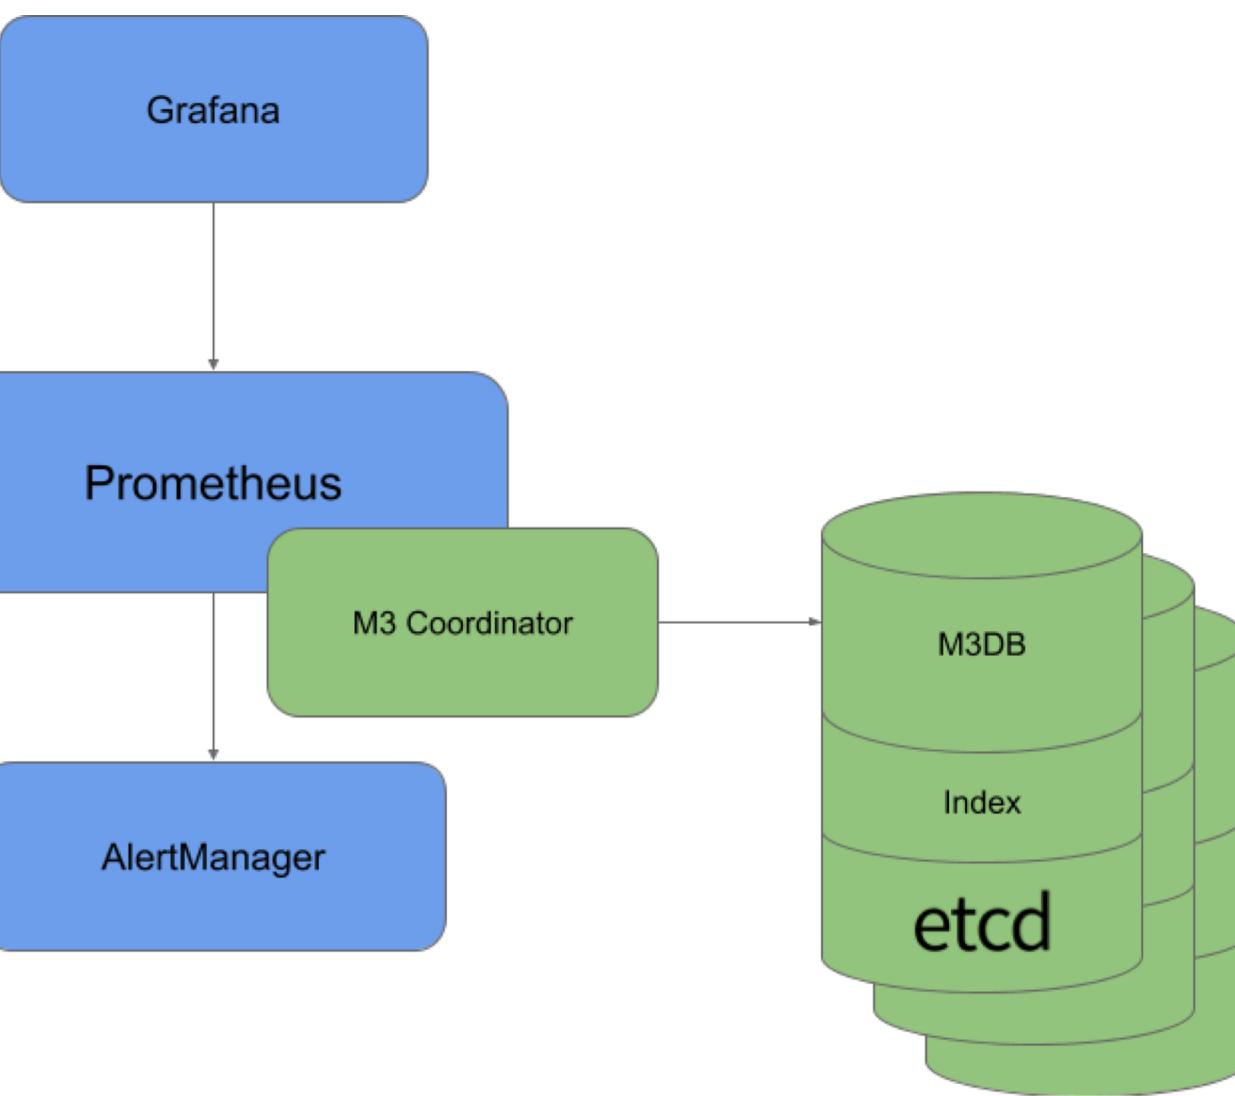




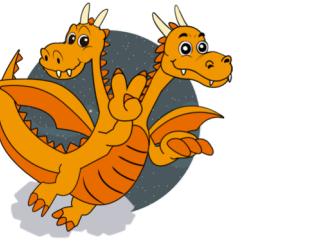
4.4 M3DB

Born to facilitate the growth of Uber's global operations

```
- name: Retain all disk_used metrics
  filter: name:disk_used* device:sd*
policies:
  - resolution: 10s
    retention: 2d
  - resolution: 1m
    retention: 30d
  - resolution: 1h
    retention: 5y
```



- Reliably houses large-scale metrics over long retention time windows
- **Downsampling of metrics**



4.5 Cortex

- Multi-tenant, horizontally scalable open source **Prometheus-as-a-Service**
- CNCF project since September 2018
- Provides a complete out-of-the-box solution for even most demanding monitoring and observability use cases.
- Hosted Cortex: Weave Cloud and Grafana Cloud

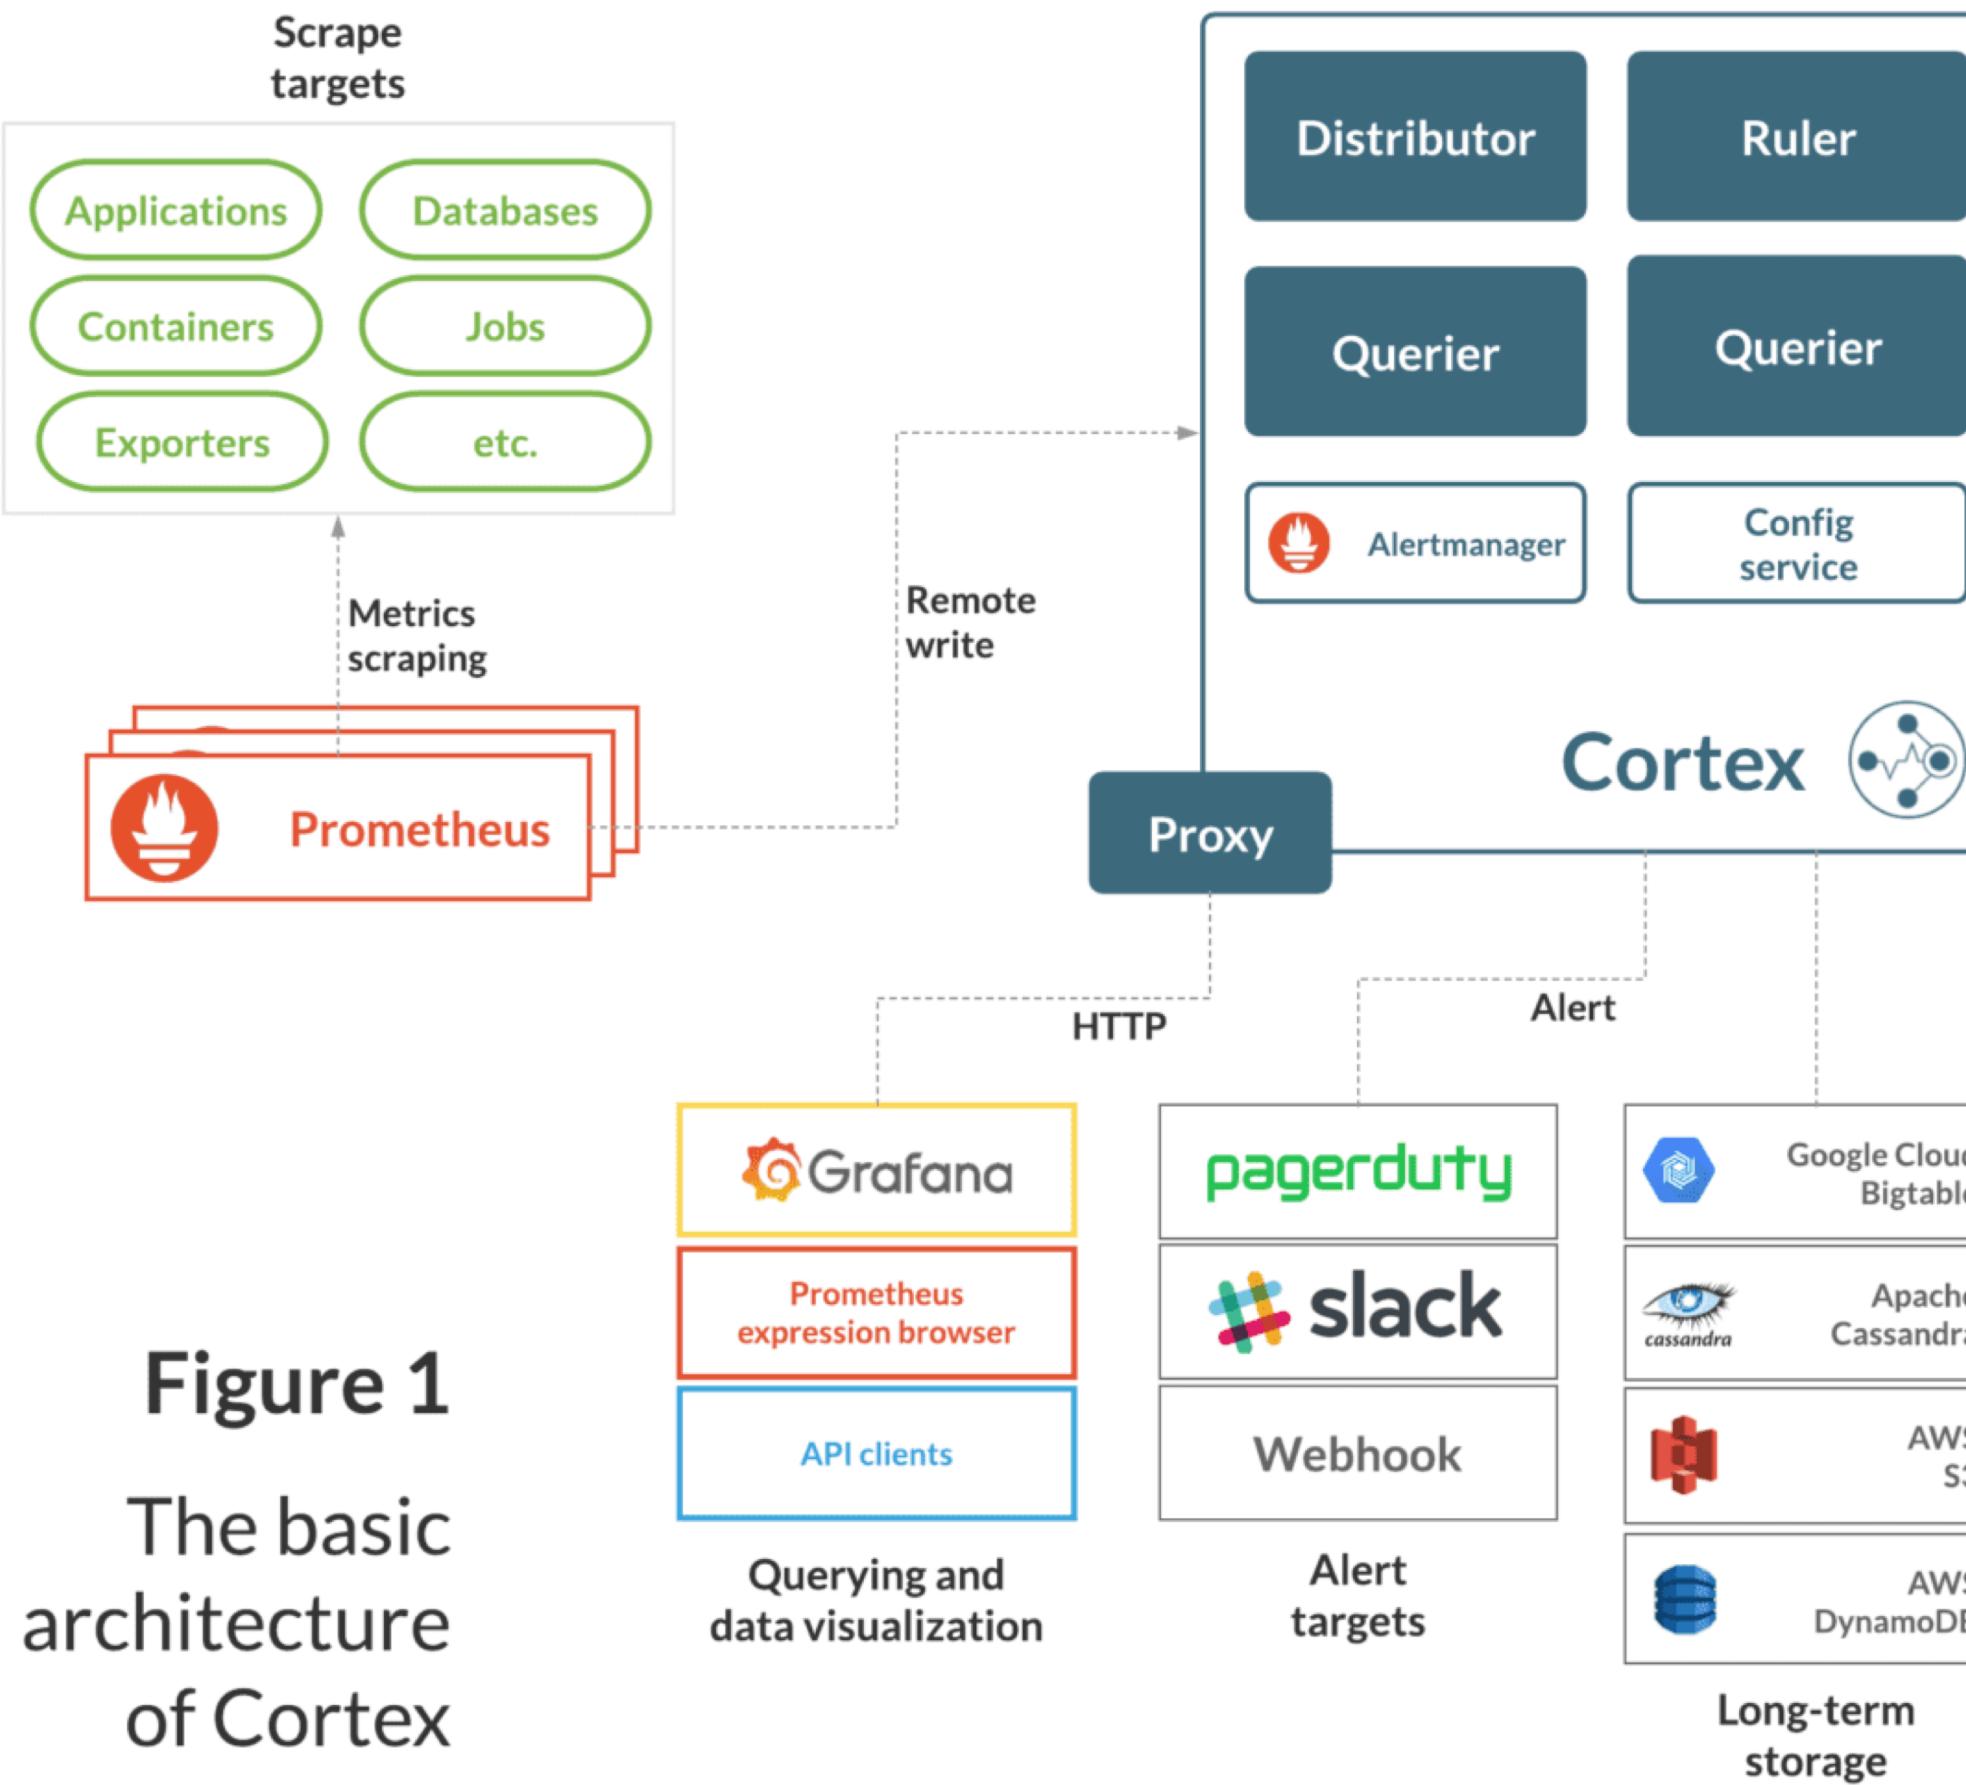
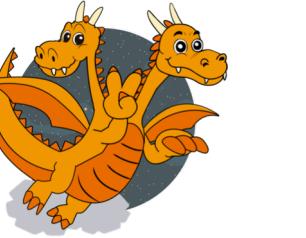


Figure 1
The basic
architecture
of Cortex





4.6 Signifai

AIOps with Prometheus

- Acquired by NewRelic
- AI and machine learning **powered correlation engine** for DevOps and Site Reliability Engineering.
- End-to-end systems analysis
- Get answers not alerts
- Automatic correlation, aggregation and prioritization of alerts.
- Be proactive not reactive



Conclusions

cloud-native monitoring: observability
RED metrics, don't fly blind
Cross team implication
Instrumentation: not as painful after all

Beatriz Martínez
@beatrizmrg



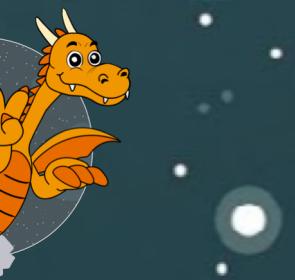
DEVOPS SPAIN II EDICIÓN

Patrocinia



Colabora





DEVOPS SPAIN II EDICIÓN



Colabora



POLITÉCNICA



Patrocina

